

Surname (COGNOME)	SOLUTION
Name	
POLIMI Personal Code	
Signature	

Politecnico di Milano, 11 February, 2025

Course on Advanced Computer Architectures

Prof. C. Silvano

EX1	(5 points)	
EX2	(3 points)	
EX3	(2 points)	
EX4	(5 points)	
Q1	(5 points)	
Q2	(5 points)	
QUIZZES	(8 points)	
TOTAL	(33 points)	

EXERCISE 1 – DEPENDENCY ANALYSIS + TOMASULO (5 points)

1. Let's consider the following assembly code containing multiple types of intra-loop dependences. Complete the following table by inserting all types of true-data-dependences, anti-dependences and output dependences for each instruction:

I#	TYPE OF INSTRUCTION	ANALYSIS OF DEPENDENCES: 1. True data dependence with I# for \$Fx 2. Anti-dependence with I# for \$Fy 3. Output-dependence with I# for \$Fz
I0	LOOP: LD \$F0, A(\$R0)	None
I1	LD \$F2, B(\$R0)	None
I2	FADD \$F4, \$F0, \$F2	True data dependence with I0 for \$F0 True data dependence with I1 for \$F2
I3	FADD \$F6, \$F4, \$F4	True data dependence with I2 for \$F4
I4	SD \$F4, C(\$R0)	True data dependence with I2 for \$F4
I5	SD \$F6, D(\$R0)	True data dependence with I3 for \$F6
I6	ADDI \$R0, \$R0, 4	Anti dependence with I0, I1, I4, I5 for \$R0
I7	BNE \$R0, \$R1, LOOP	True data dependence with I6 for \$R0

2. Let's consider the previous assembly code to be executed on a CPU with dynamic scheduling based on TOMASULO algorithm with all cache HITS, a single Common Data Bus and:

- 2 RESERV. STATIONS (RS1, RS2) with 2 LOAD/STORE units (LDU1, LDU2) with latency 6
- 2 RESERVATION STATION (RS3, RS4) with 2 FP unit12 (FPU1, FPU2) with latency 2
- 1 RESERVATION STATION (RS5) with 1 INT_ALU/BR unit (ALU1) with latency 1

Please complete the following table:

INSTRUCTION	ISSUE	START EXEC	WRITE RESULT	Hazards Type	RSi	UNIT
LOOP: LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW \$F0 RAW \$F2	RS3	ALU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW \$F4	RS4	ALU2
SD \$F4, C(\$R0)	9	13	19	STRUCT RS1 RAW \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10	16	22	RAW \$F6	RS2	LDU2
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15	16	STRUCT RS4 (RAW \$R0)	RS5	ALU2

(*) Only hazards that have caused the introduction of some stalls are reported in the table. Other dependencies are already solved.

(**) In Tomasulo, WAW and WAR hazards are already solved by Register Renaming in ISSUE stage.

Calculate the **CPI** = **CPI** = # clock cycles / IC = 22 / 8 = 2,75

EXERCISE 2 – VLIW SCHEDULING (3 points)

Let's consider the following LOOP code:

```

LOOP: LD F1, A(R1)
      LD F2, A(R2)
      LD F3, A(R3)
      FADD F1, F1, F2
      FADD F2, F2, F3
      FMUL F1, F1, F2
      FADD F3, F3, F3
      ADDUI R2, R1, 8
      ADDUI R3, R1, 8
      SD F1, B(R1)
      ADDUI R1, R1, 4
      BNE R1, R6, LOOP
    
```

Given a 3-issue VLIW machine with **fully pipelined functional units**:

- 1 Memory Unit with 3 cycles latency
- 1 FP ALU with 2 cycles latency
- 1 Integer ALU with 1 cycle latency to next Int/FP & 2 cycle latency to next Branch

The branch is completed with 1 cycle delay slot (branch solved in ID stage). **No branch prediction.**

In the Register File, it is possible to read and write at the same address at the same clock cycle.

Considering one iteration of the loop, complete the following table by using the **list-based scheduling** (do NOT introduce any software pipelining, loop unrolling and modifications to loop indexes) on the 4-issue VLIW machine including the **BRANCH DELAY SLOT**. Please do not write in NOPs.

	Memory Unit 1	Floating Point Unit 1	Integer Unit
C1	LD F1, A(R1)		
C2	LD F2, A(R2)		ADDUI R2, R1, 8
C3	LD F3, A(R3)		ADDUI R3, R1, 8
C4			
C5		FADD F1, F1, F2	
C6		FADD F2, F2, F3	
C7		FADD F3, F3, F3	
C8		FMUL F1, F1, F2	
C9			
C10	SD F1, B(R1)		ADDUI R1, R1, 4
C11			
C12			BNE R1, R2, LOOP
C13			Br. delay slot
C14			
C15			

1. How long is the critical path?

13 cycles (There is NO branch prediction, so the branch delay slot cannot be used for next iteration)

2. How much is the code efficiency?

$$\text{Code_eff} = \text{IC} / (\# \text{ cycles} * \# \text{ issues}) = 12 / (13 * 3) = 12 / 39 = 0.3$$

EXERCISE 2 – MESI PROTOCOL (2 points)

Let's consider the following access patterns on a **4-processor** system with a direct-mapped, write-back cache with one cache block per processor and a two-cache block memory.

Assume the **MESI protocol** is used with **write-back** caches, **write-allocate**, and **write-invalidate** of other caches.

Please **COMPLETE** the following table:

Cycle	After Operation	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0	Excl (1)	Invalid	Shared (0)	Shared (0)	Yes	Yes
4	P1: Write Bl. 0	Excl (1)	Mod (0)	Invalid	Invalid	No	Yes
5	P0: Write Bl. 1	Mod (1)	Mod (0)	Invalid	Invalid	No	No
6	P3: Read Bl. 1	Shared(1)	Mod (0)	Invalid	Shared (1)	No	Yes

EXERCISE 4 on REORDER BUFFER (5 points)

Let's consider the following assembly loop where registers **\$R1** and **\$R2** are initialized at 0 and 40 respectively:

```
L0:    LD $F2, 0 ($R1)
L1:    ADDD $F4, $F2, $F2
L2:    SD $F4, 0 ($R1)
L3:    ADDI $R1, $R1, 4
L4:    BNE $R1, $R2, L0 # branch predicted as taken
```

- How many loop iterations? *10 iterations*
- How many instructions per iteration? *5 instructions*
- How many control instructions vs datapath instructions? *2 vs 3*

1. Write the unrolled version of the loop with unrolling factor 2 by using Register Renaming:

L0:	LD \$F2, 0 (\$R1)
L1:	ADDD \$F4, \$F2, \$F2
L2:	SD \$F4, 0 (\$R1)
L3:	LD \$F6, 4 (\$R1)
L4:	ADDD \$F8, \$F6, \$F6
L5:	SD \$F8, 4 (\$R1)
L6:	ADDI \$R1, \$R1, 8
L7:	BNE \$R1, \$R2, LOOP

- How many loop iterations? *5 iterations*
- How many instructions per iteration? *8 instructions*
- How many control instructions vs datapath instructions? *2 vs 6*

2. Execute the unrolled version of the loop by the **Speculative Tomasulo** architecture with a **10-entry ROB** and:

- 4 Load Buffers (Load1, Load2, Load3, Load4);
- 4 FP Reservation Stations (FP1, FP2, FP3, FP4)
- 2 Integer Reservation Stations (Int1, Int2)

Complete the ROB and the Rename Table until the ROB becomes **full** while the first instruction is still in execution due to a cache miss (*):

ROB Table

ROB#	Instruction	Dest.	Resource Allocation	Ready /Status	Spec.	
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1	No, exec.(*)	No	HEAD
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1	No, issued	No	
ROB2	L2: SD \$F4, 0 (\$R1)	Mem	--	No, issued	No	
ROB3	L3: LD \$F6, 4 (\$R1)	\$F6	Load2	No, issued	No	
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2	No, issued	No	
ROB5	L5: SD \$F8, 4 (\$R1)	Mem	--	No, issued	No	
ROB6	L6: ADDI \$R1, \$R1, 8	\$R1	Int1	No, issued	No	
ROB7	L7: BNE \$R1, \$R2, LOOP	--	Int2	No, issued	No	
ROB8	L0: LD \$F2, 0 (\$R1)	\$F2	Load3	No, issued	Yes	
ROB9	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP3	No, issued	Yes	

Rename Table:

Reg#	ROB#
\$F0	--
\$F2	ROB0, ROB8
\$F4	ROB1, ROB9
\$F6	ROB3
\$F8	ROB4

QUESTION 1: Instruction-Level and Thread-Level Parallelism (5 points)

Modern processors exploit Instruction Level Parallelism and Thread Level Parallelism.

Answer to the following questions:

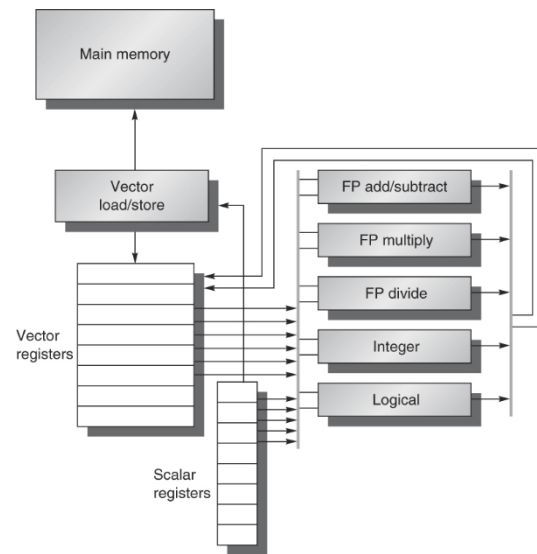
	Instruction Level Parallelism (ILP)	Thread Level Parallelism (TLP)
<i>Explain the main concepts for each approach.</i>		
<i>Which type of technique can be applied in superscalar processors to combine both ILP and TLP?</i>		

<i>Explain what type of instruction scheduling is used in superscalar processors to combine both ILP and TLP?</i>	
<i>What are the main hardware modifications required to a generic superscalar processor to support both ILP and TLP?</i>	

QUESTION 2: VECTOR PROCESSORS (5 points)

Consider a vector processor architecture, as VMIPS shown in the figure.

- Present the main concepts;
- Present the main advantages of vector execution with respect to scalar execution, detailing the features of the vector architecture that provide the advantage.



QUIZZES (8 points)

Question 1 (format Multiple Choice – Single answer)

Let's consider the following code executed by a Vector Processor with:

- Vector Register File composed of 32 vectors of 8 elements per 64 bits/element;
- Scalar FP Register File composed of 32 registers of 64 bits;
- One Load/Store Vector Unit with operation chaining and memory bandwidth 64 bits;
- One ADD/SUB Vector Unit with operation chaining;
- One MUL/DIV Vector Unit with operation chaining.

```
L.V V1, RA      # Load vector from memory address RA into V1
L.V V3, RB      # Load vector from memory address RB into V3
MULVS.D V1, V1, F0 # FP multiply vector V1 to scalar F0
ADDVV.D V2, V1, V1 # FP add vectors V1 and V1
MULVS.D V2, V2, F0 # FP multiply vector V2 to scalar F0
ADDVV.D V3, V2, V3 # FP add vectors V2 and V3
S.V V1, RX      # Store vector V3 into memory address RX
S.V V2, RY      # Store vector V3 into memory address RY
S.V V3, RZ      # Store vector V3 into memory address RZ
```

How many convoys? How many clock cycles to execute the code?

(SINGLE ANSWER)

1 point

Answer 1: 2 convoys; 16 clock cycles

Answer 2: 3 convoys; 24 clock cycles

Answer 3: 4 convoys; 32 clock cycles

Answer 4: 5 convoys; 40 clock cycles **(TRUE)**

Answer 5: 6 convoys; 48 clock cycles

Motivate your answer by completing the following table:

1 point

	Load/Store Vector Unit	Add/Sub Vector Unit	Mul/Div Vector Unit
1^ convoy	L.V V1, RA;	ADDVV.D V2, V1, V1	MULVS.D V1, V1, F0
2^ convoy	L.V V3, RB	ADDVV.D V3, V2, V3	MULVS.D V2, V2, F0
3^ convoy	S.V V1, RX		
4^ convoy	S.V V2, RY		
5^ convoy	S.V V3, RZ		

Another possible solution is:

	Load/Store Vector Unit	Add/Sub Vector Unit	Mul/Div Vector Unit
1^ convoy	L.V V1, RA;	ADDVV.D V2, V1, V1	MULVS.D V1, V1, F0
2^ convoy	S.V V1, RX		
3^ convoy	L.V V3, RY	ADDVV.D V3, V2, V3	MULVS.D V2, V2, F0
4^ convoy	S.V V2, RY		
5^ convoy	S.V V3, RZ		

Question 2 (format True/False)

In VLIW architectures, the compiler can detect parallelism only in basic blocks of the code;

(format True/False)

1 point

Answer 1: True / False **(False)**

Motivate your answer:

1 point

Question 3 (format True/False)

A 4-issue VLIW processor requires 4 Program Counters to load the necessary instructions in the 4 parallel lanes

(format True/False)

1 point

Answer 1: True / False **(False)**

Motivate your answer:

1 point

Question 4 (format Multiple Choice – Single answer)

In the *Speculative Tomasulo Architecture*, what type of hardware block is used to undo speculative instructions in case of a mispredicted branch?

(SINGLE ANSWER)

1 point

Answer 1: Reorder Buffer; **(TRUE)**

Answer 2: Instruction Dispatcher;

Answer 3: Store Buffers;

Answer 4: Reservation Stations;

Answer 5: Load Buffers;

Motivate your answer by explaining what happens on a mispredicted branch:

1 point
