

| | |
|-----------------------------|--|
| Surname (COGNOME) | |
| Name | |
| POLIMI Personal Code | |
| Signature | |

Politecnico di Milano, 11 February, 2025

Course on Advanced Computer Architectures

Prof. C. Silvano

| | | |
|----------------|--------------------|--|
| EX1 | (5 points) | |
| EX2 | (3 points) | |
| EX3 | (2 points) | |
| EX4 | (5 points) | |
| Q1 | (5 points) | |
| Q2 | (5 points) | |
| QUIZZES | (8 points) | |
| TOTAL | (33 points) | |

EXERCISE 1 – DEPENDENCY ANALYSIS + TOMASULO (5 points)

1. Let's consider the following assembly code containing multiple types of intra-loop dependences. Complete the following table by inserting all types of true-data-dependences, anti-dependences and output dependences for each instruction:

| I# | TYPE OF INSTRUCTION | ANALYSIS OF DEPENDENCES: 1. True data dependence with I# for \$Fx 2. Anti-dependence with I# for \$Fy 3. Output-dependence with I# for \$Fz |
|----|------------------------|--|
| I0 | LOOP: LD \$F0, A(\$R0) | None |
| I1 | LD \$F2, B(\$R0) | None |
| I2 | FADD \$F4, \$F0, \$F2 | |
| I3 | FADD \$F6, \$F4, \$F4 | |
| I4 | SD \$F4, C(\$R0) | |
| I5 | SD \$F6, D(\$R0) | |
| I6 | ADDI \$R0, \$R0, 4 | |
| I7 | BNE \$R0, \$R1, LOOP | |

2. Let's consider the previous assembly code to be executed on a CPU with dynamic scheduling based on TOMASULO algorithm with all cache HITS, a single Common Data Bus and:

- 2 RESERV. STATIONS (RS1, RS2) with 2 LOAD/STORE units (LDU1, LDU2) with latency 6
- 2 RESERVATION STATION (RS3, RS4) with 2 FP unit12 (FPU1, FPU2) with latency 2
- 1 RESERVATION STATION (RS5) with 1 INT_ALU/BR unit (ALU1) with latency 1

Please complete the following table:

| | | | | | | |
|------------------------|---|---|---|------|-----|------|
| LOOP: LD \$F0, A(\$R0) | 1 | 2 | 8 | None | RS1 | LDU1 |
| LD \$F2, B(\$R0) | 2 | 3 | 9 | None | RS2 | LDU2 |
| FADD \$F4, \$F0, \$F2 | | | | | | |
| FADD \$F6, \$F4, \$F4 | | | | | | |
| SD \$F4, C(\$R0) | | | | | | |
| SD \$F6, D(\$R0) | | | | | | |
| ADDI \$R0, \$R0, 4 | | | | | | |
| BNE \$R0, \$R1, LOOP | | | | | | |

Calculate the CPI = _____

EXERCISE 2 – VLIW SCHEDULING (3 points)

Let's consider the following LOOP code:

```

LOOP: LD F1, A(R1)
      LD F2, A(R2)
      LD F3, A(R3)
      FADD F1, F1, F2
      FADD F2, F2, F3
      FMUL F1, F1, F2
      FADD F3, F3, F3
      ADDUI R2, R1, 8
      ADDUI R3, R1, 8
      SD F1, B(R1)
      ADDUI R1, R1, 4
      BNE R1, R6, LOOP
    
```

Given a 3-issue VLIW machine with fully pipelined functional units:

- 1 Memory Unit with 3 cycles latency
- 1 FP ALU with 2 cycles latency
- 1 Integer ALU with 1 cycle latency to next Int/FP & 2 cycle latency to next Branch

The branch is completed with 1 cycle delay slot (branch solved in ID stage). **No branch prediction.**

In the Register File, it is possible to read and write at the same address at the same clock cycle.

Considering one iteration of the loop, complete the following table by using the **list-based scheduling** (do NOT introduce any software pipelining, loop unrolling and modifications to loop indexes) on the 4-issue VLIW machine including the **BRANCH DELAY SLOT**. Please do not write in NOPs.

| | Memory Unit 1 | Floating Point Unit 1 | Integer Unit |
|------------|---------------|-----------------------|--------------|
| C1 | LD F1, A(R1) | | |
| C2 | | | |
| C3 | | | |
| C4 | | | |
| C5 | | | |
| C6 | | | |
| C7 | | | |
| C8 | | | |
| C9 | | | |
| C10 | | | |
| C11 | | | |
| C12 | | | |
| C13 | | | |
| C14 | | | |
| C15 | | | |

1. How long is the critical path? _____

2. How much is the code efficiency? _____

EXERCISE 2 – MESI PROTOCOL (2 points)

Let's consider the following access patterns on a **4-processor** system with a direct-mapped, write-back cache with one cache block per processor and a two-cache block memory.

Assume the **MESI protocol** is used with **write-back** caches, **write-allocate**, and **write-invalidate** of other caches.

1. Please COMPLETE the following table:

| Cycle | After Operation | P0 cache block state | P1 cache block state | P2 cache block state | P3 cache block state | Mem. at bl. 0 up to date? | Mem. at bl. 1 up to date? |
|-------|-----------------|----------------------|----------------------|----------------------|----------------------|---------------------------|---------------------------|
| 1 | P0: Read Bl. 1 | Excl (1) | Invalid | Invalid | Invalid | Yes | Yes |
| 2 | P2: Read Bl. 0 | Excl (1) | Invalid | Excl (0) | Invalid | Yes | Yes |
| 3 | P3: Read Bl. 0 | | | | | | |
| 4 | P1: Write Bl. 0 | | | | | | |
| 5 | P0: Write Bl. 1 | | | | | | |
| 6 | P3: Read Bl. 1 | | | | | | |

EXERCISE 4 on REORDER BUFFER (5 points)

Let's consider the following assembly loop where registers **\$R1** and **\$R2** are initialized at 0 and 40 respectively:

```

L0:    LD $F2, 0 ($R1)
L1:    ADDD $F4, $F2, $F2
L2:    SD $F4, 0 ($R1)
L3:    ADDI $R1, $R1, 4
L4:    BNE $R1, $R2, L0 # branch predicted as taken
    
```

- How many loop iterations? _____
- How many instructions per iteration? _____
- How many control vs datapath instructions? _____

1. Write the unrolled version of the loop with unrolling factor 2 by using Register Renaming:

| | |
|-----|-----------------------|
| L0: | LD \$F2, 0 (\$R1) |
| L1: | ADDD \$F4, \$F2, \$F2 |
| L2: | |
| | |
| | |
| | |
| | |
| | |
| | |
| | |

- How many loop iterations? _____
- How many instructions per iteration? _____
- How many control vs datapath instructions _____

2. Execute the unrolled version of the loop by the **Speculative Tomasulo** architecture with a **10-entry ROB** and:

- 4 Load Buffers (Load1, Load2, Load3, Load4);
- 4 FP Reservation Stations (FP1, FP2, FP3, FP4)
- 2 Integer Reservation Stations (Int1, Int2)

Complete the ROB and the Rename Table until the ROB becomes **full** while the first instruction is still in execution due to a cache miss (*):

ROB Table

| ROB# | Instruction | Dest. | Resource Allocation | Ready /Status | Spec. | |
|------|---------------------------|-------|---------------------|---------------|-------|------|
| ROB0 | L0: LD \$F2, 0 (\$R1) | \$F2 | Load1 | No, exec. (*) | No | HEAD |
| ROB1 | L1: ADDD \$F4, \$F2, \$F2 | \$F4 | FP1 | No, issued | No | |
| ROB2 | | | | | | |
| ROB3 | | | | | | |
| ROB4 | | | | | | |
| ROB5 | | | | | | |
| ROB6 | | | | | | |
| ROB7 | | | | | | |
| ROB8 | | | | | | |
| ROB9 | | | | | | |

Rename Table:

| Reg# | ROB# |
|------|------|
| \$F0 | -- |
| \$F2 | |
| \$F4 | |
| \$F6 | |
| \$F8 | |

QUESTION 1: Instruction-Level and Thread-Level Parallelism (5 points)

Modern processors exploit Instruction Level Parallelism and Thread Level Parallelism.

Answer to the following questions:


| | Instruction Level Parallelism (ILP) | Thread Level Parallelism (TLP) |
|--|--|-----------------------------------|
| <i>Explain the main concepts for each approach.</i> | | |
| <i>Which type of technique can be applied in superscalar processors to combine both ILP and TLP?</i> | | |

| | |
|---|--|
| <p><i>Explain what type of instruction scheduling is used in superscalar processors to combine both ILP and TLP?</i></p> | |
| <p><i>What are the main hardware modifications required for a generic superscalar processor to support both ILP and TLP?</i></p> | |

QUESTION 2: VECTOR PROCESSORS (5 points)

Consider a vector processor architecture, as VMIPS shown in the figure.

- Present the main concepts;
- Present the main advantages of vector execution with respect to scalar execution, detailing the features of the vector architecture that provide the advantage.



QUIZZES (8 points)

Question 1 (format Multiple Choice – Single answer)

Let's consider the following code executed by a Vector Processor with:

- Vector Register File composed of 32 vectors of 8 elements per 64 bits/element;
- Scalar FP Register File composed of 32 registers of 64 bits;
- One Load/Store Vector Unit with operation chaining and memory bandwidth 64 bits;
- One ADD/SUB Vector Unit with operation chaining;
- One MUL/DIV Vector Unit with operation chaining.

L.V V1, RA # Load vector from memory address RA into V1
L.V V3, RB # Load vector from memory address RB into V3
MULVS.D V1, V1, F0 # FP multiply vector V1 to scalar F0
ADDVV.D V2, V1, V1 # FP add vectors V1 and V1
MULVS.D V2, V2, F0 # FP multiply vector V2 to scalar F0
ADDVV.D V3, V2, V3 # FP add vectors V2 and V3
S.V V1, RX # Store vector V3 into memory address RX
S.V V2, RY # Store vector V3 into memory address RY
S.V V3, RZ # Store vector V3 into memory address RZ

How many convoys? How many clock cycles to execute the code?

(SINGLE ANSWER)

1 point

Answer 1: 2 convoys; 16 clock cycles

Answer 2: 3 convoys; 24 clock cycles

Answer 3: 4 convoys; 32 clock cycles

Answer 4: 5 convoys; 40 clock cycles

Answer 5: 6 convoys; 48 clock cycles

Motivate your answer by completing the following table:

1 point

| | Load/Store Vector Unit | Add/Sub Vector Unit | Mul/Div Vector Unit |
|-----------|------------------------|---------------------|---------------------|
| 1^ convoy | L.V V1, RA; | | |
| 2^ convoy | | | |
| 3^ convoy | | | |
| 4^ convoy | | | |
| 5^ convoy | | | |
| 6^ convoy | | | |

Question 2 (format True/False)

In VLIW architectures, the compiler can detect parallelism only in basic blocks of the code;

(format True/False)

1 point

Answer 1: True / False

Motivate your answer:

1 point

Question 3 (format True/False)

A 4-issue VLIW processor requires 4 Program Counters to load the necessary instructions in the 4 parallel lanes

(format True/False)

1 point

Answer 1: True / False

Motivate your answer:

1 point

Question 4 (format Multiple Choice – Single answer)

*In the **Speculative Tomasulo Architecture**, what type of hardware block is used to undo speculative instructions in case of a mispredicted branch?*

(SINGLE ANSWER)

1 point

Answer 1: Reorder Buffer;

Answer 2: Instruction Dispatcher;

Answer 3: Store Buffers;

Answer 4: Reservation Stations;

Answer 5: Load Buffers;

Motivate your answer by explaining what happens in the case of a mispredicted branch:

1 point
