

Course on: “Advanced Computer Architectures”

The problem of control hazards



Prof. Cristina Silvano
Politecnico di Milano
email: cristina.silvano@polimi.it




Conditional Branch Instructions

- Examples of conditional branches for RISC-V processor:
beq (*branch on equal*) and **bne** (*branch on not equal*)
beq rs1, rs2, L1 # go to L1 if (rs1 == rs2)
bne rs1, rs2, L1 # go to L1 if (rs1 != rs2)
- The **Branch Target Address** is the address where to branch.
- If the branch condition is satisfied => the branch is **taken** and the **Branch Target Address** is stored in the Program Counter (PC).
- Otherwise => the branch is **not taken** or **untaken** and the instruction stream is executed sequentially with the next instruction address (**PC + 4**).



Execution of branches for 5-stage RISC-V pipeline

beq rx, ry, L1



IF: Instruction Fetch and PC increment

ID: Instruction Decode and Registers Read (**rx** and **ry**)

EX: Compare registers (**rx** and **ry**) in the ALU to derive the
Branch Outcome: Taken / Not Taken.

Computation of **Branch Target Address: (PC+offset)**

ME: The Branch Outcome is used to decide the next PC:

PC <= (PC+4) if Branch Not Taken

PC <= (PC+offset) if Branch Taken




```
beq  rx,ry,L1
```

- **Branch Outcome (T/NT)** and **BTA** are ready at the end of EX stage.
- Branches are solved when **PC** is updated at the end of ME stage.



Problem of Control Hazards in Pipelining (1)

- To feed the pipeline we need to fetch a new instruction at each clock cycle, but the branch decision (to change or not change the PC) is not yet ready!



- The problem to choose the correct instruction to be fetched after a branch is called *Control Hazard* or *Branch Hazard*

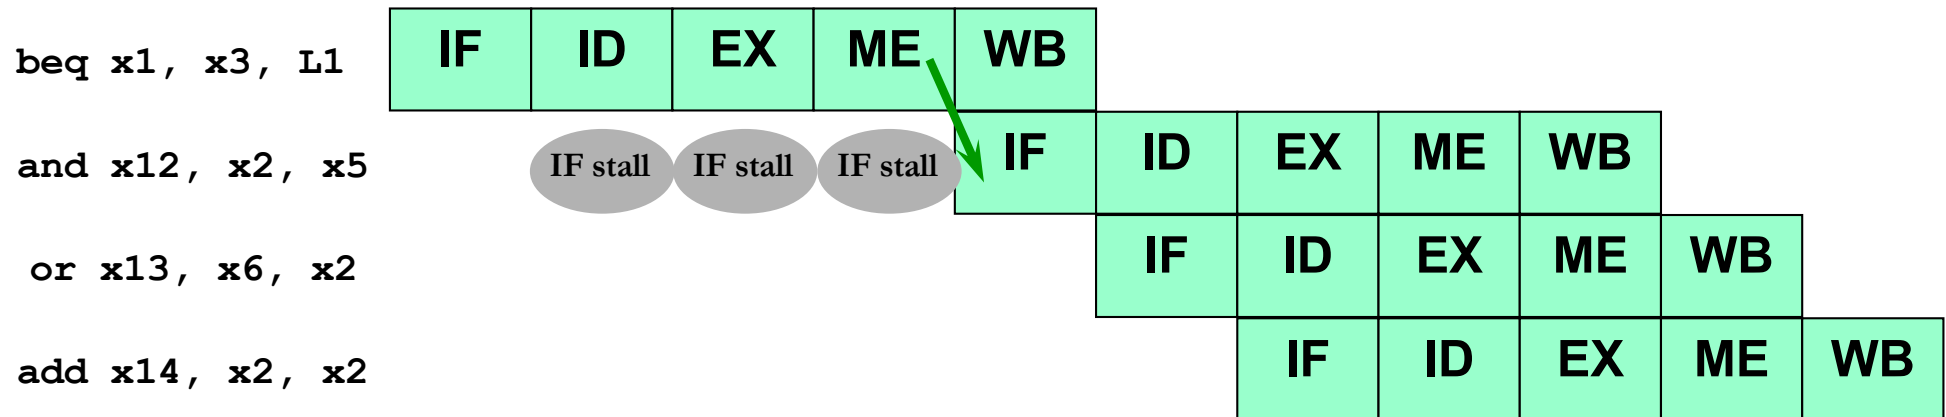


Problem of Control Hazards in Pipelining (2)

- **Control hazards:** Attempt to take a decision on the next instruction to fetch in the pipeline *before* the branch condition is evaluated and PC updated.
- **Control hazards** arise from the pipelining of conditional branches and other *jump* instructions changing the PC.
- **Control hazards** reduce the performance from the ideal speedup gained by the pipelining because *it is needed to stall the pipeline until branch resolution.*



Branch Stalls: Conservative Solution



- **Conservative assumption:** *Stalling until resolution* at the end of the ME stage of the branch.
- **Performance reduction:** Each branch costs a penalty of *3 stalls* to decide and fetch the correct instruction flow in the pipeline.



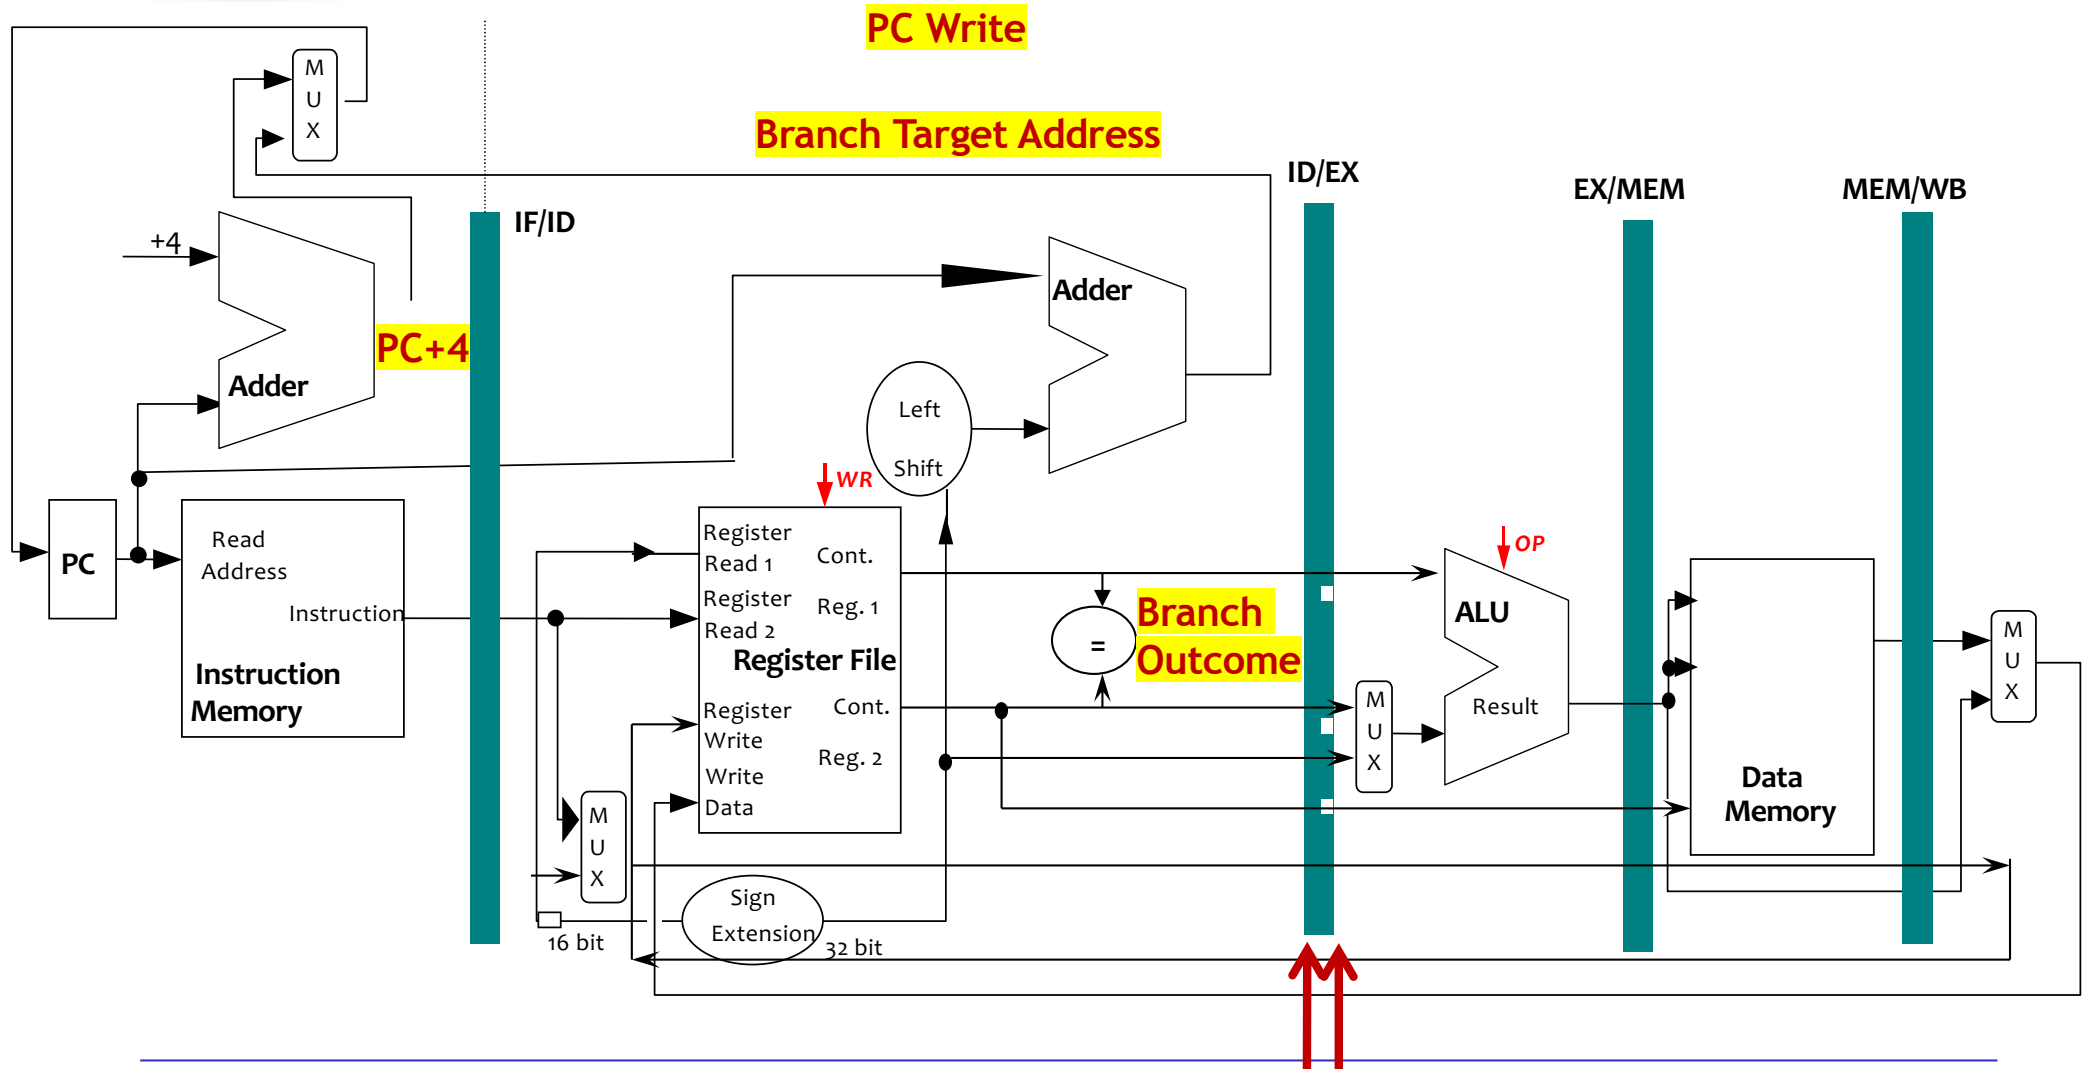
Early Evaluation of branch in ID stage

- **Idea:** To improve performance in case of branch hazards, we need to add more hardware resources to:
 1. Compare registers to derive the **Branch Outcome**
 2. Compute the **Branch Target Address**
 3. Update the PC register*as soon as possible in the pipeline.*

*==> RISC-V optimized pipeline anticipates the 1,2,3 steps **during the ID stage** ==> Both BO and BTA are known at the end of ID stage.*

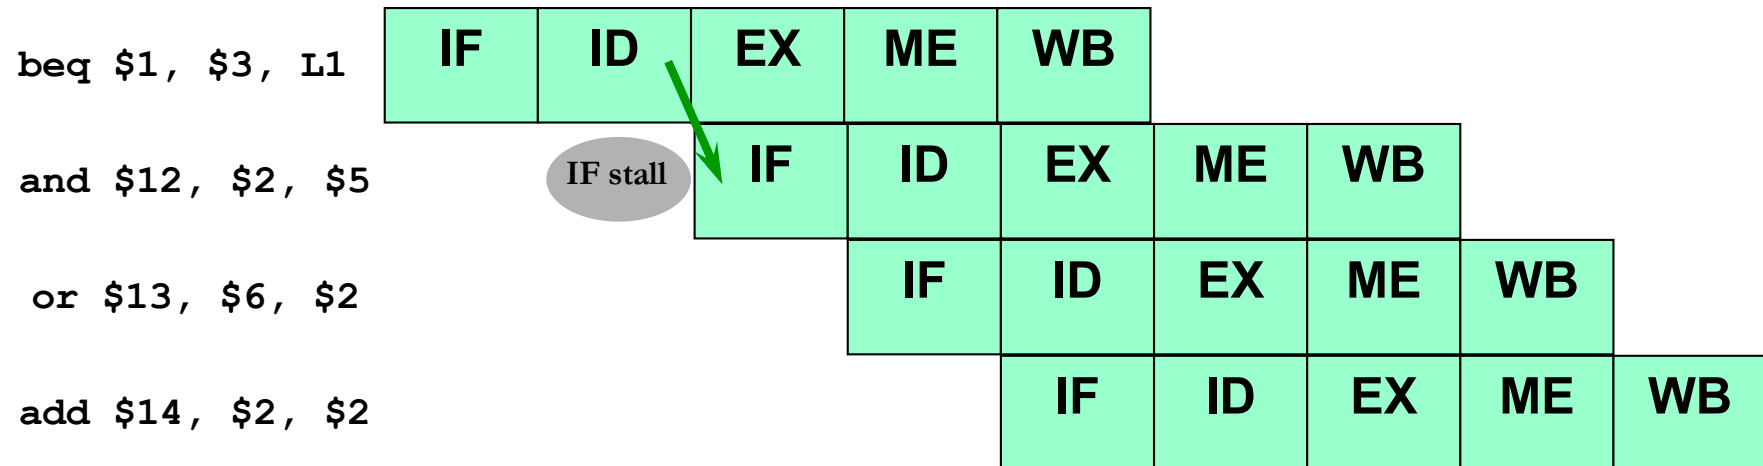


Early Evaluation of branch in ID stage (2)





Early Evaluation of branch in ID stage (3)



- **Conservative assumption:** Stalling until resolution at the end of the ID stage (when the BO and BTA are known) to decide which instruction to fetch => Now each branch costs **one stall penalty** to fetch the correct instruction.

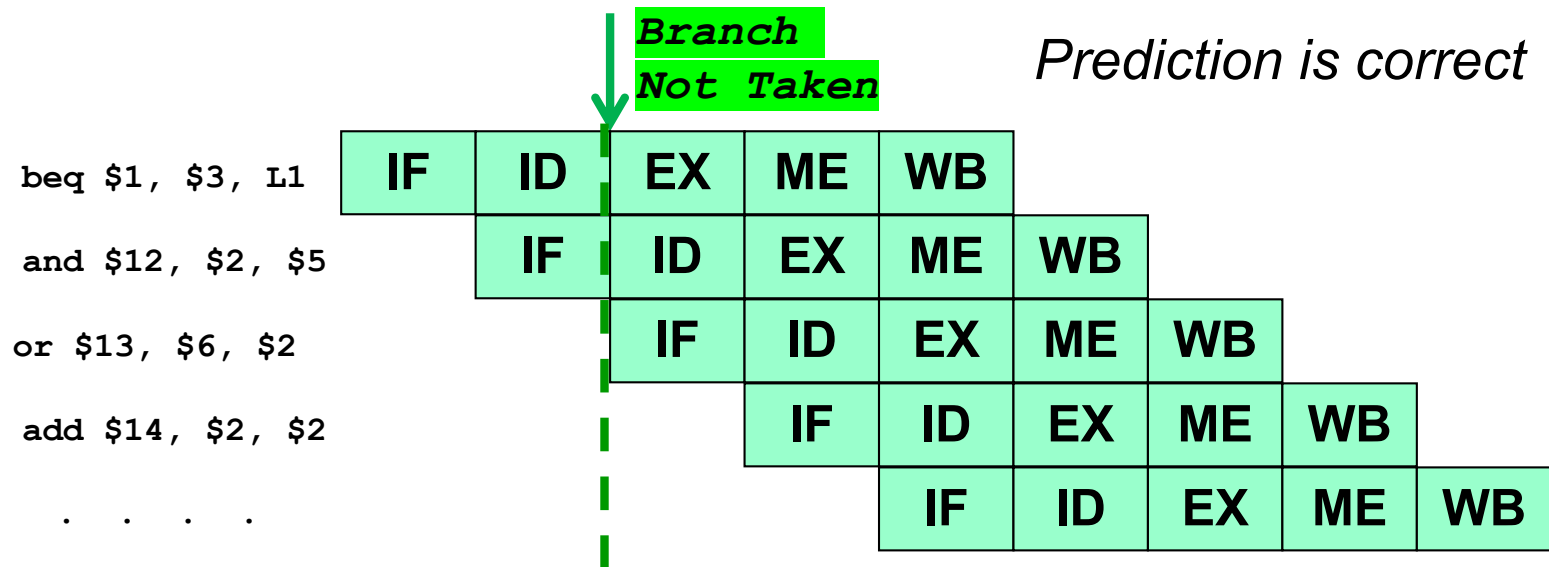


Branch Hazards: Possible Solution

- Possible solution: Let's assume the *branch as not taken* and continue to execute sequentially.
- *Only if* the branch outcome will be *taken* => flush the next instruction in the pipeline.
- Be careful: We cannot assume the *branch as taken* because we don't know yet the branch target address!
- This solution introduces the idea of *branch prediction*.



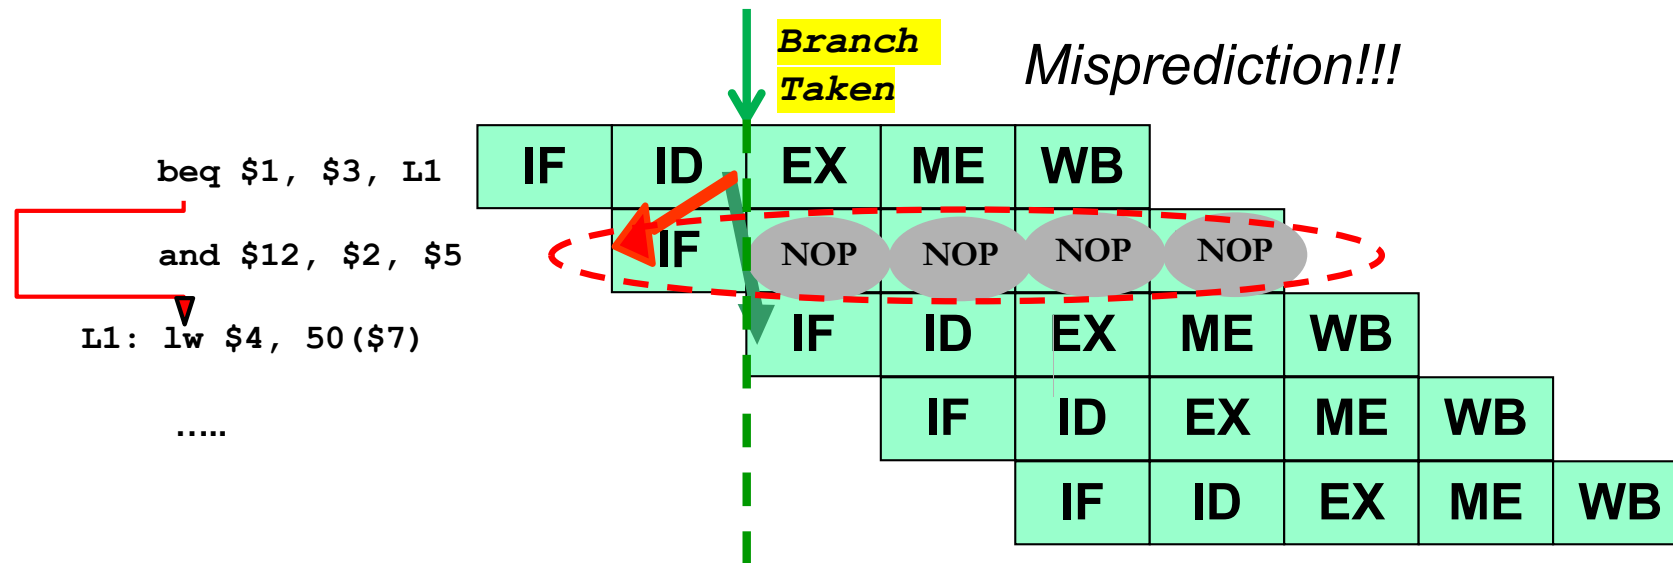
Let's predict the branch as not taken (1)



- Let's assume the branch as *not taken* :
 1. If the branch outcome will be *not taken*, the pipeline execution of sequential instructions is fine!
 2. *What happens if the branch will be taken?*



Let's predict the branch as not taken (1)



- If the branch outcome will be **taken**, it will be necessary to:
 1. **Flush** only **one** instructions before writing its results;
 2. Fetch next instruction at the **Branch Target Address (L1)**
- => In this case, each branch costs **one stall penalty**.



Summary

- With the branch decision is anticipated at ID stage:
=> reduction of the cost associated to each branch (**branch penalty**):
 - We need **only one-clock-cycle IF stall** after each branch
 - Or a **flush** of only **one** instruction following the branch
- One-cycle-delay for every branch still yields a performance loss of about 10% to 30% depending on the branch frequency:

*Stall Cycles per Instruction due to Branches =
Branch Frequency x Branch Penalty*

- Let's introduce in the next chapter some **branch prediction techniques** to deal with this performance loss.