

Course on: “Advanced Computer Architectures”

---


# Introduction to Instruction Level Parallelism

---



Prof. Cristina Silvano  
Politecnico di Milano  
email: [cristina.silvano@polimi.it](mailto:cristina.silvano@polimi.it)

# Sequential vs. Pipelining Execution



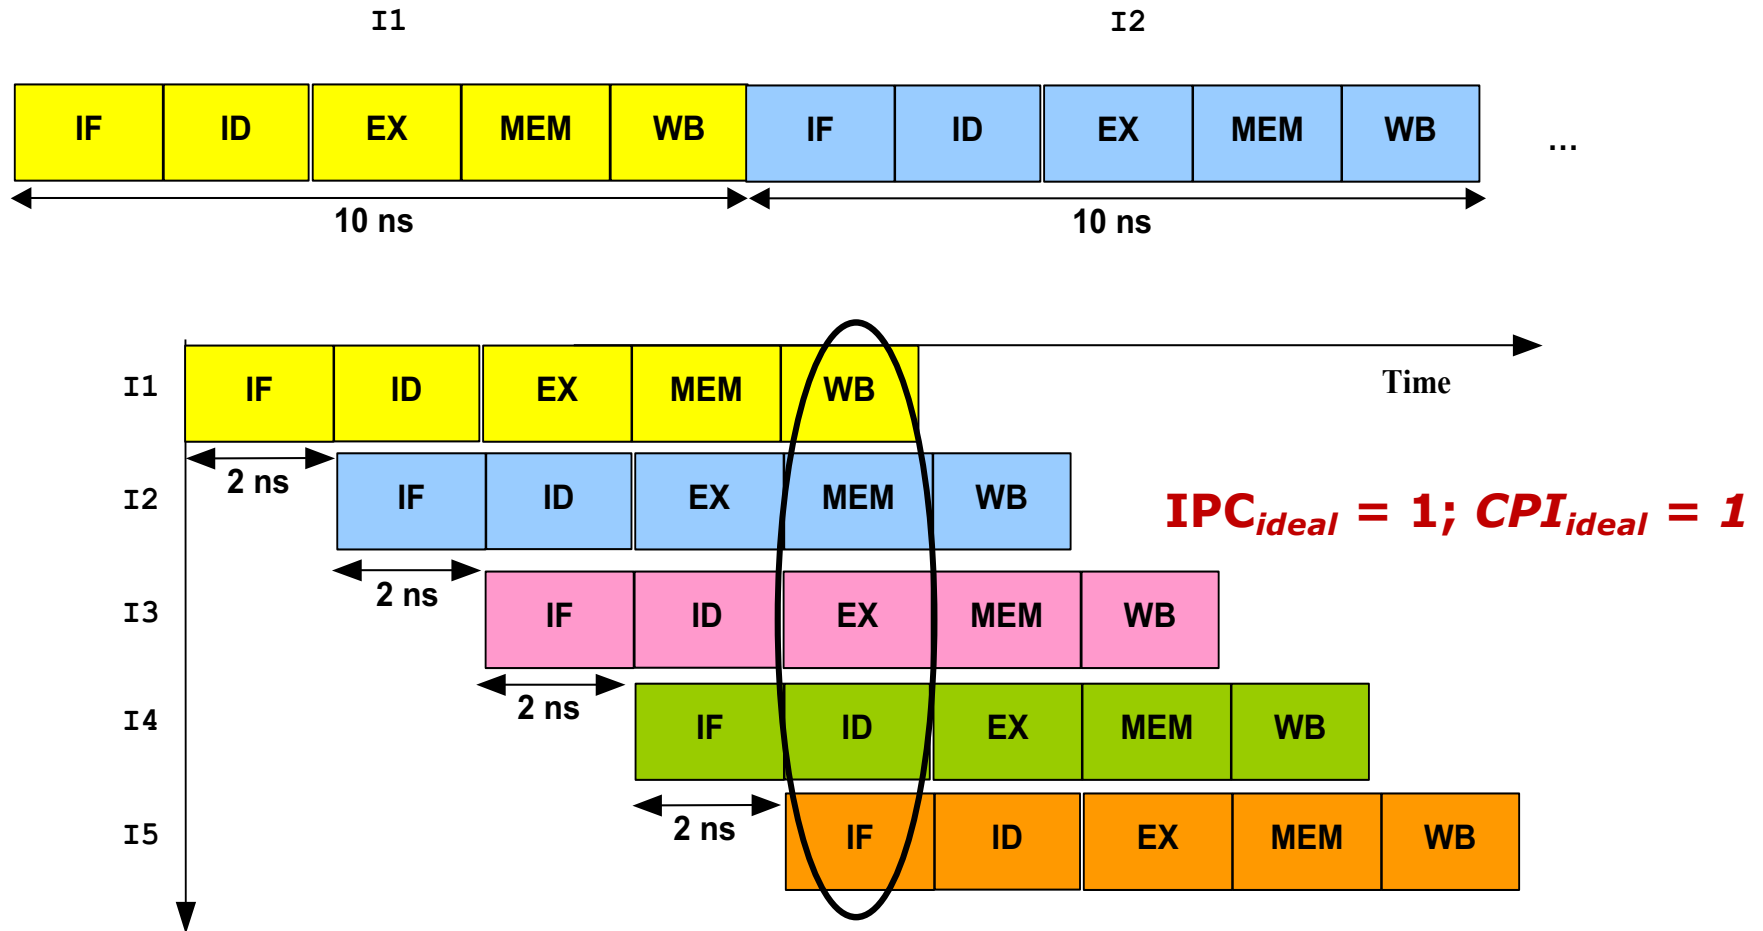
***5-stage pipeline => 5 different instructions overlapped***

# Pipelining Execution

---

- Pipelining overlaps the execution of instructions  
=> it exploits the "***Instruction Level Parallelism***"
- When exploiting ILP, the goal is to maximize performance in terms of throughput: **Instructions per clock (IPC)**  
... and minimize the  **$CPI = 1 / IPC$**
- *Remind that **pipelining** improves instruction throughput, but not the latency of the single instruction!*

# Best case in pipelining execution



# The problem of dependences

- Determining **dependences** among instructions is critical to define the amount of parallelism existing in a program.
- If two instructions are **dependent** to each other, they cannot be executed in parallel: they must be executed in order or only partially overlapped.
- There are **three** different types of dependences in a code:
  1. **True Data Dependences:** an instruction j is dependent on a data produced by a previous instruction i;
  - ➡ 2. **Name Dependences:** two instructions use the same register or memory location;
  3. **Control Dependences:** they impose the ordering of instructions

# Name Dependences

---

- A **name dependence** occurs when two instructions use the same register or memory location (called **name**), but there is **no flow of data** between the instructions associated with that name
- Name dependences are **not** true data dependences, since there is no value (no data flow) being transmitted between the two instructions => this is just a **register reuse!**
- There are **two types** of name dependences:
  1. **Anti-dependences**
  2. **Output Dependences**

# Two types of name dependences

- Let's consider **Ii** that precedes instruction **Ij** in program order:

- Anti-dependence:**

When **Ij** writes a register or memory location that instruction **Ii** reads => *it can generate a **Write After Read – WAR hazard***.  
Original instructions ordering must be preserved to ensure that **Ii** reads the previous value:

Ii:  $r_3 \leftarrow (r_1) \text{ op } (r_2)$

Ij:  $r_1 \leftarrow (r_4) \text{ op } (r_5)$

- Output Dependence:**

When **Ii** and **Ij** write the same register or memory location => *it can generate a **Write After Write – WAW hazard***.  
Original instructions ordering must be preserved to ensure that the value finally written corresponds to **Ij**.

Ii:  $r_3 \leftarrow (r_1) \text{ op } (r_2)$

Ij:  $r_3 \leftarrow (r_6) \text{ op } (r_7)$

# How to solve name dependences?

- Name dependences are **not** true data dependences, since there is no value passed (no data flow) between instructions.
- Name dependences are due to **register reuse**.
- **Register Renaming:** If the register used could be changed, then the instructions do not conflict anymore.
- Examples:

Ii:  $r_3 \leftarrow (r_1) \text{ op } (r_2)$   
Ij:  $r_1 \leftarrow (r_4) \text{ op } (r_5) \Rightarrow r_4 \leftarrow (r_4) \text{ op } (r_5)$

Ii:  $r_3 \leftarrow (r_1) \text{ op } (r_2)$   
Ij:  $r_3 \leftarrow (r_6) \text{ op } (r_7) \Rightarrow r_4 \leftarrow (r_6) \text{ op } (r_7)$



# Register Renaming

---

- **Register Renaming** can be more easily done, if there are enough registers available in the ISA.
- **Register Renaming** can be done either statically by the compiler or dynamically by the hardware.
- Dependences through memory locations are more difficult to detect ("**memory disambiguation**" problem), since two addresses may refer to the same location but can look different.

# Data Dependences and Hazards

- A data/name dependence can potentially generate a data hazard (**RAW, WAW, or WAR**), but the actual hazard and the number of stalls to eliminate the hazards are a property of the pipeline architecture.
  - **RAW hazards** correspond to **true data dependences**
  - **WAR hazards** correspond to **anti-dependences**
  - **WAW hazards** correspond to **output dependences**
- **Dependences** are a property of the program, while **hazards** are a property of the pipeline architecture.

# Summary on Data Dependences & Hazards

---

Consider executing a sequence of

$$r_k \leftarrow (r_i) \text{ op } (r_j)$$

## True Data-dependence

$r_3 \leftarrow (r_1) \text{ op } (r_2)$   
 $r_5 \leftarrow (r_3) \text{ op } (r_4)$       Read-after-Write (RAW) hazard



## Anti-dependence

$r_3 \leftarrow (r_1) \text{ op } (r_2)$   
 $r_1 \leftarrow (r_4) \text{ op } (r_5)$       Write-after-Read (WAR) hazard



## Output-dependence

$r_3 \leftarrow (r_1) \text{ op } (r_2)$   
 $r_3 \leftarrow (r_6) \text{ op } (r_7)$       Write-after-Write (WAW) hazard



# Summary of Control Dependences

- A **control dependence** determines the ordering of instructions and it is preserved by two properties:
  - Instructions execution in **program order** to ensure that an instruction that occurs before a branch is executed before the branch.
  - Detection of control hazards to ensure that an instruction (that is control-dependent on a branch) is not executed until the branch direction is known.
- Although preserving control dependence is a simple way to preserve program order, **control dependence is not the critical property** that must be preserved (as seen when we've studied scheduling techniques to fill in the branch delay slot).

# Program Properties

---

- **Two properties** are critical to preserve program correctness (and normally preserved by maintaining both data and control dependences during scheduling):
  1. **Data flow:** Actual flow of data values among instructions that produces the correct results and consumes them.
  2. **Exception behavior:** Preserving exception behavior means that any changes in the ordering of instruction execution must not change how exceptions are raised in the program.

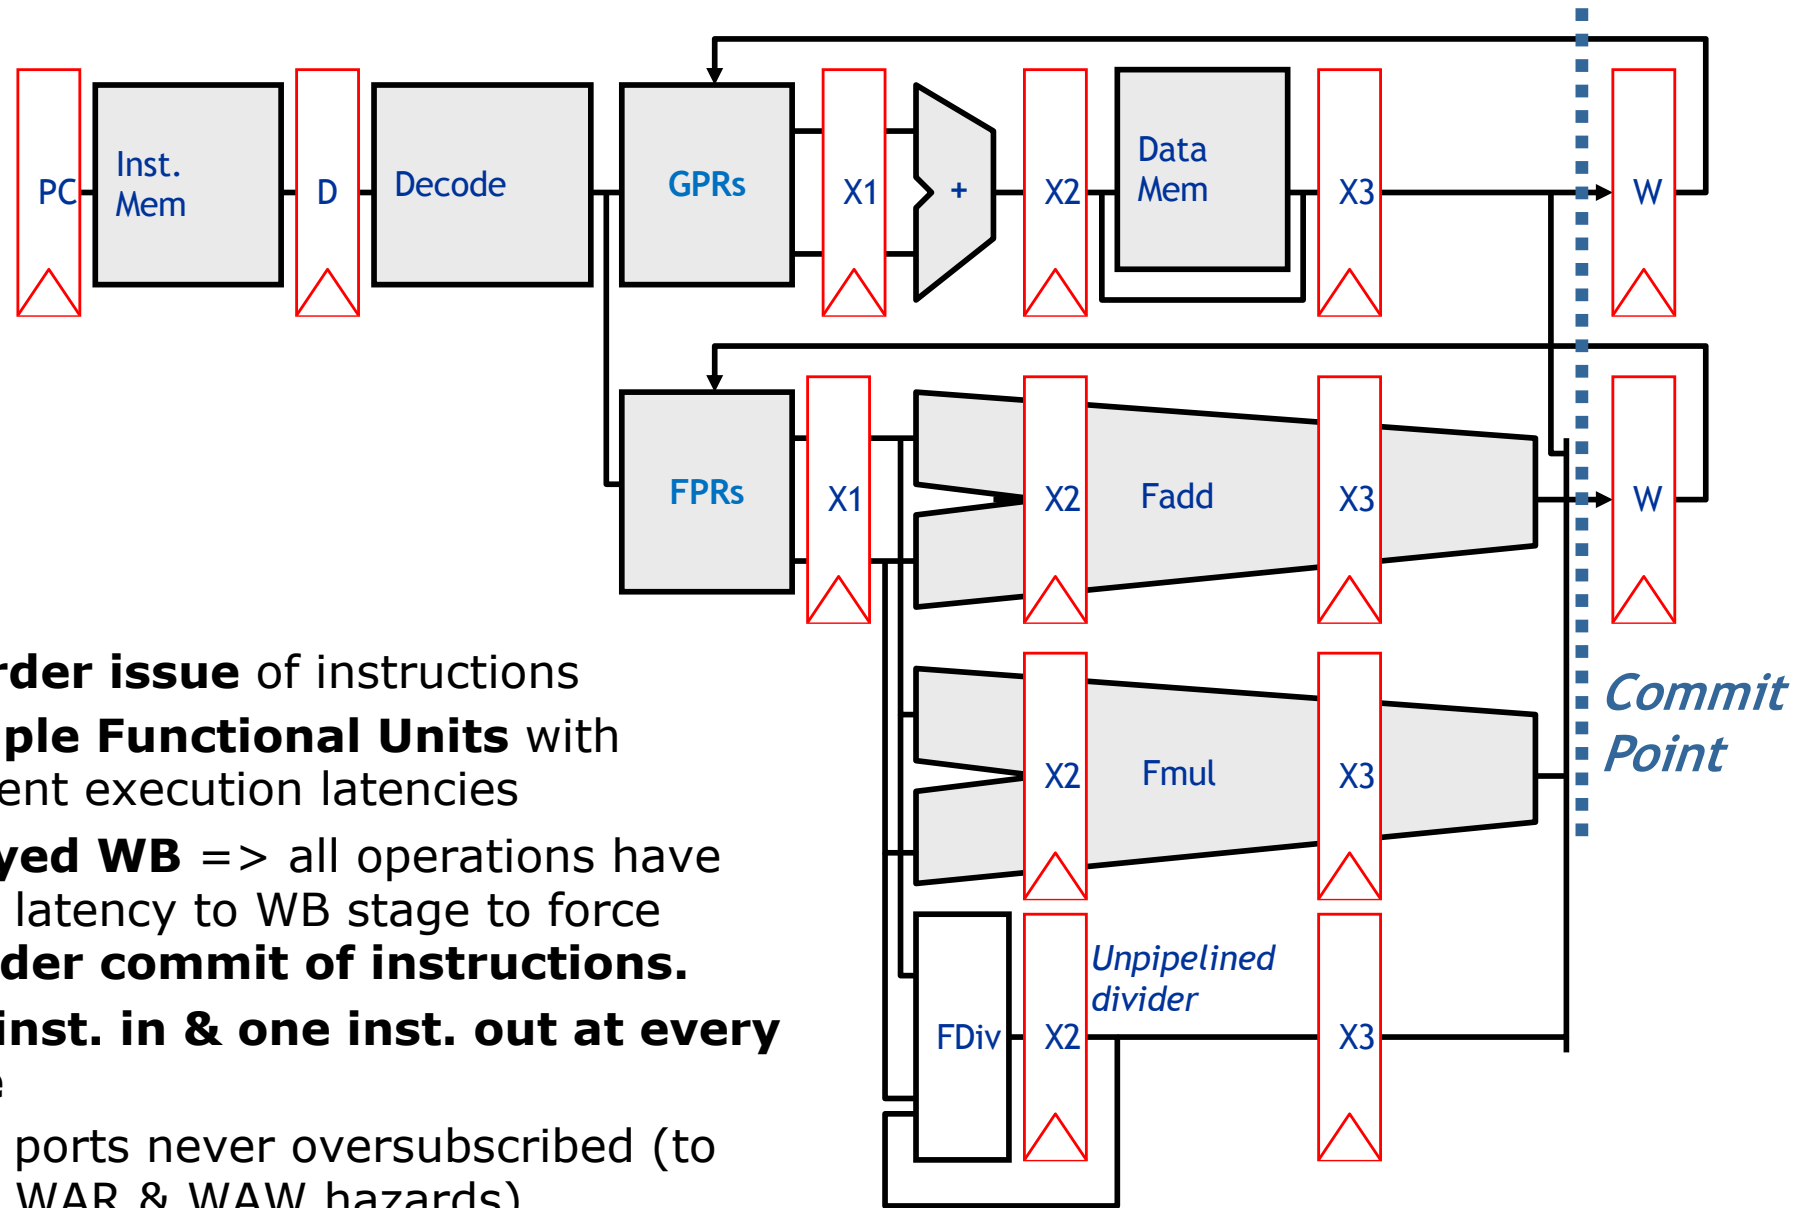
---

# **MULTI-CYCLE PIPELINING & DYNAMIC SCHEDULING**

# Multi-cycle Pipeline: Basic Assumptions

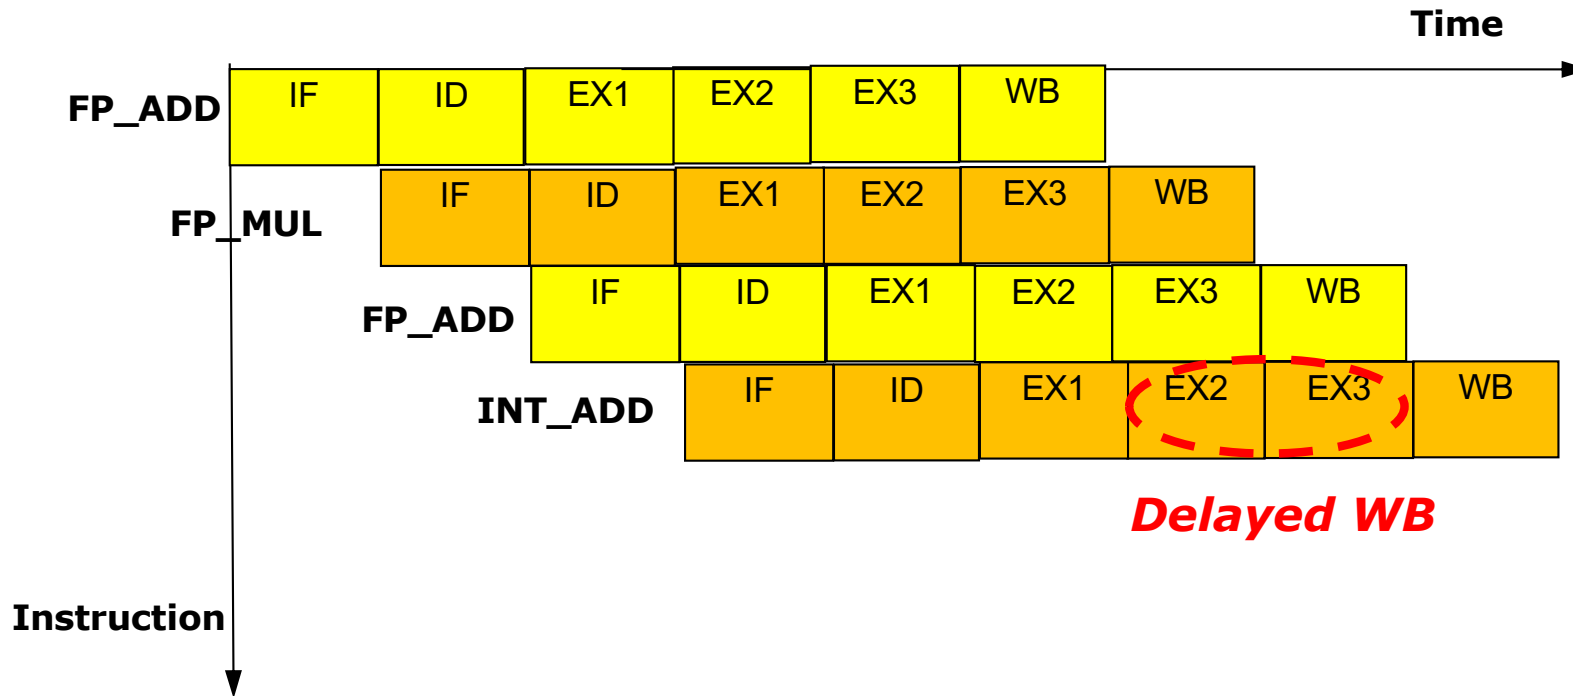
- We consider ***single-issue*** processors (one instruction issued per clock cycle).
- Instructions are then ***issued in-order.***
- **Execution stage** might require ***multiple cycles latency,*** depending on the operation type (i.e., multiply operations are typically longer than add/sub operations)
- **Memory stages** might require ***multiple cycles*** access time due to instruction and data cache misses.

# Multi-cycle In-Order Pipeline



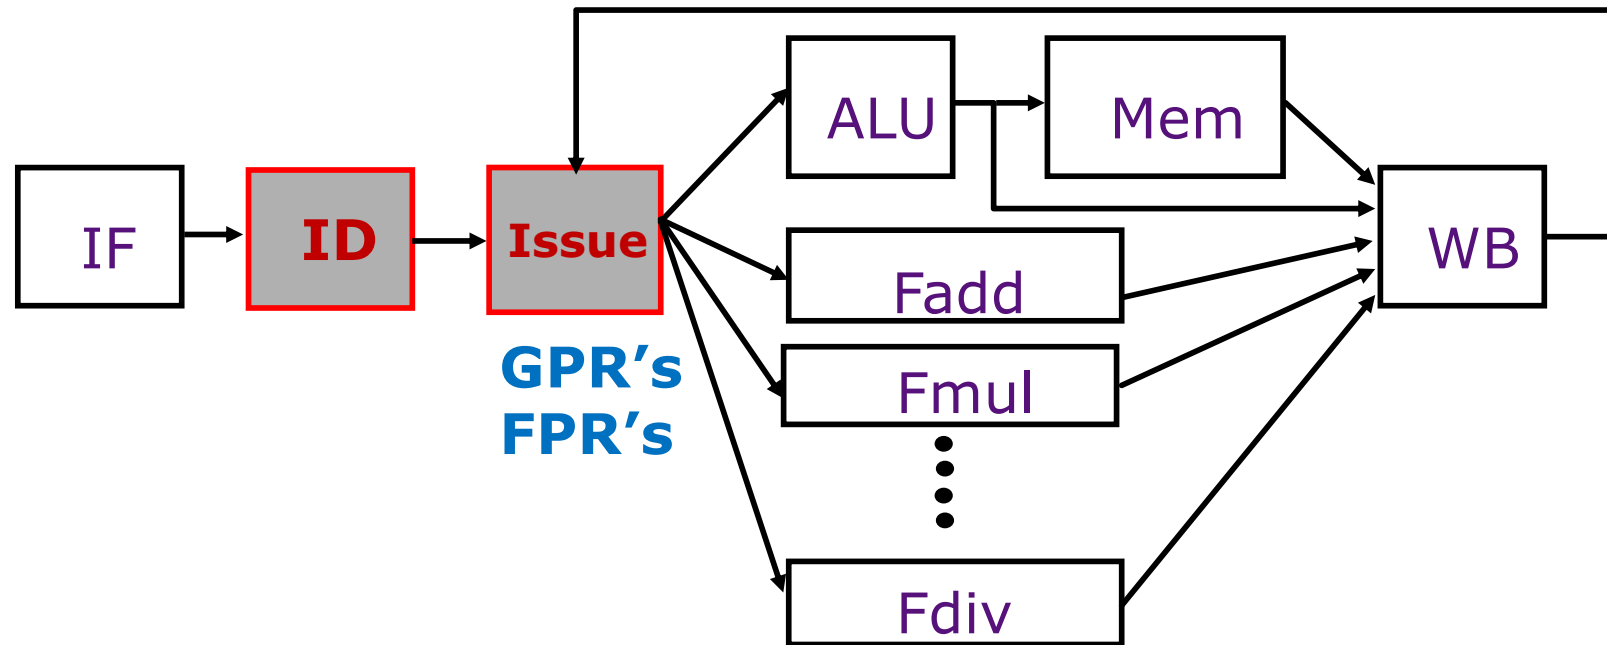


# Multi-cycle In-order Pipeline



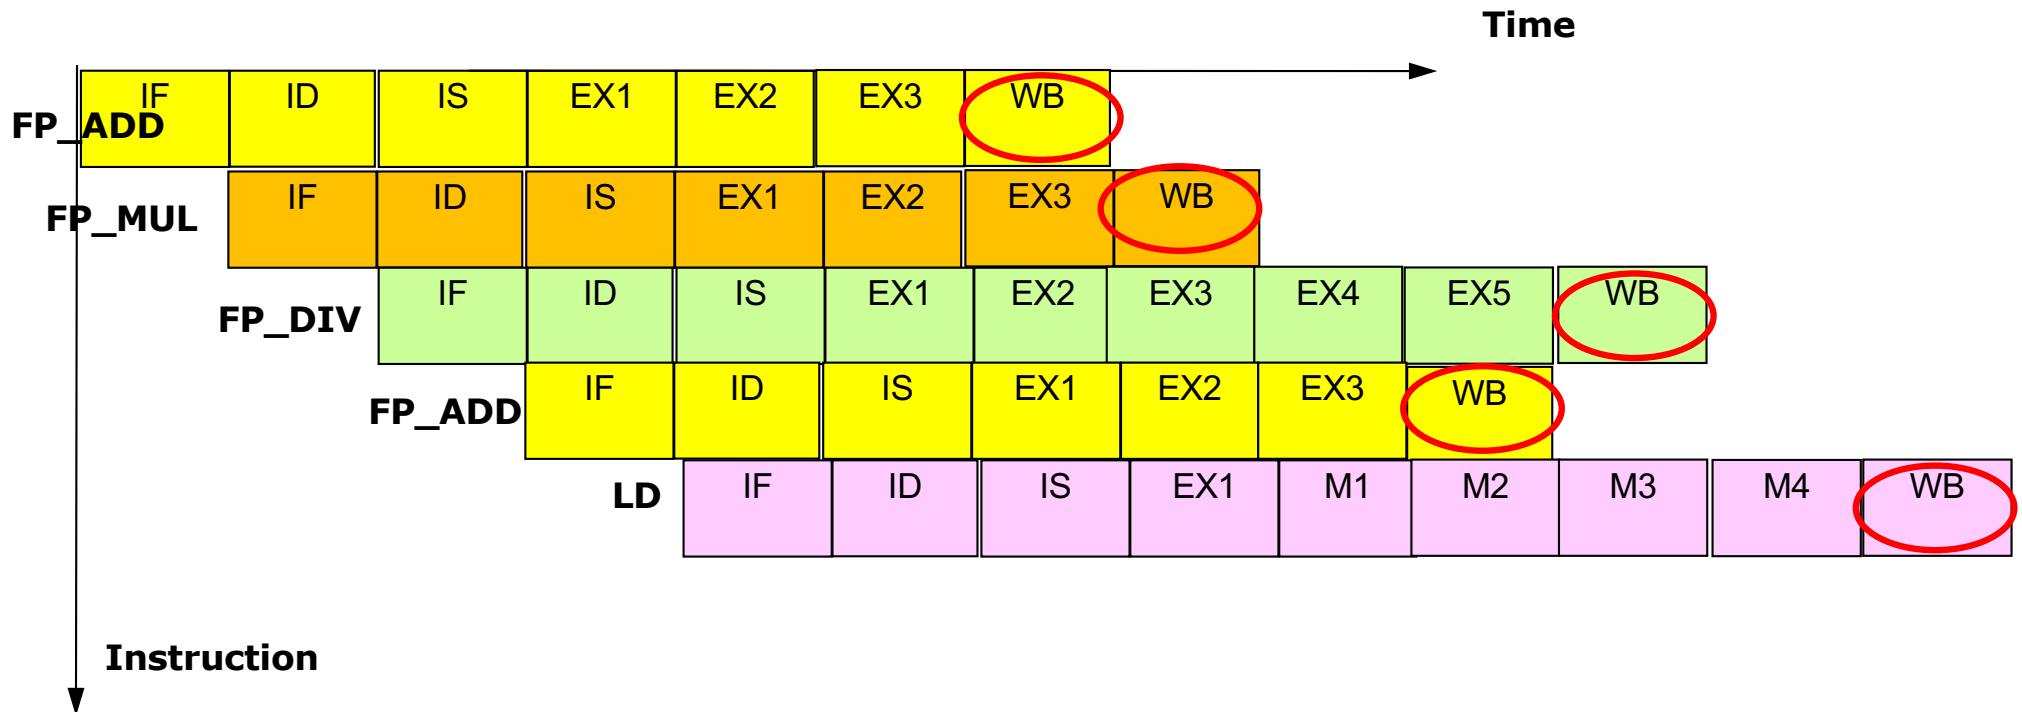
- **In-order issue & In-order commit** of instructions.
- This avoids the generation of **WAR** & **WAW hazards** and preserves the **precise exception model**.

# Multi-cycle Out-of-order Pipeline



- ID stage split in **2 stages**: Instr. Decode (**ID**) & Register Read (**Issue**);
- **Multiple** functional units with variable latency;
- Long latency **multi-cycle floating-point instructions**;
- Memory systems with variable access time: **Multi-cycle memory accesses** due to data cache misses (unpredictable statically);
- **No more commit point** => **Out-of-order commit**: Need to check for WAR & WAW hazards and imprecise exception

# Multi-cycle Out-of-order Pipeline



- **In-order issue** of instructions
- **Out-of-order execution & out-of-order commit** of instructions
- Need to check the generation of **WAR** & **WAW** hazards and **imprecise exceptions**.

# Another Key Idea: Dynamic Scheduling


- **Problem:** Hazards due to true data dependences that cannot be solved by forwarding cause **stalls** of the pipeline => no new instructions are fetched nor issued even if they are not data dependent;
- **Solution: Allow data independent instructions behind a stall to proceed**
  - HW manages dynamically the instruction execution to reduce stalls: an instruction execution begins as soon as their operands are available.
- **This generates out-of-order execution and completion (commit)**
- *Dynamic Scheduling can be applied to a single-issue scalar pipeline, but we'll see that it can be applied also to multiple-issue pipelines...*

# Example of dynamic scheduling

DIVD **F0**, F2, F4      # exec. takes many cycles  
ADDD F10, **F0**, F8      # RAW **F0**  
SUBD F12, F8, F14

- **ADDD** stalls for RAW hazard on **F0** (waiting many clock cycles for DIVD commit).
- **SUBD** would stall even if not data dependent on anything in the pipeline.
- **BASIC IDEA:** to enable **SUBD** to proceed  
=> this generates *out-of-order execution*

# Example of dynamic scheduling (cont.)



# Imprecise exceptions

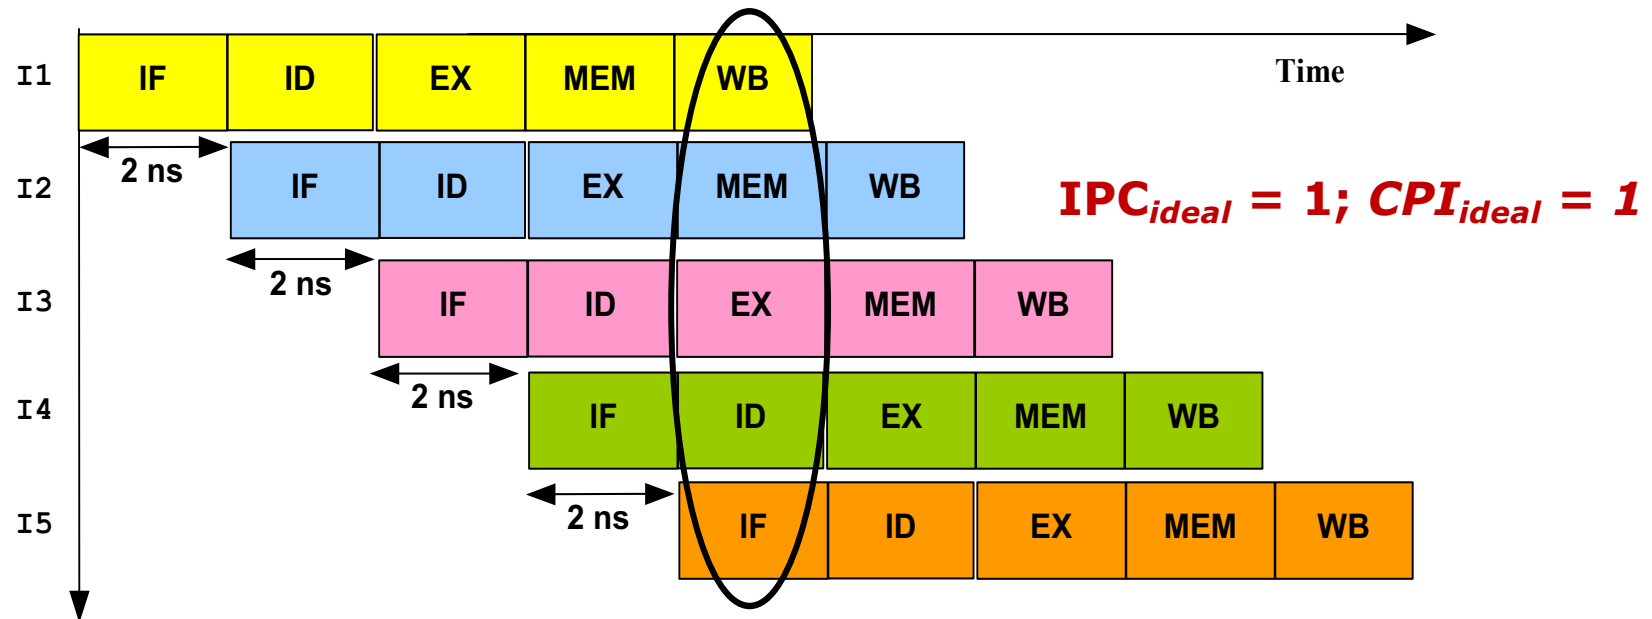
- An exception is **imprecise** if the processor state when an exception is raised does not look exactly as if the instructions were executed in-order.
- Imprecise exceptions can occur in **out-of-order** because:
  - The pipeline may have **already** completed instructions that are **later** in program order than the instruction causing the exception
  - The pipeline may have **not yet** completed some instructions that are **earlier** in program order than the instruction causing the exception
- Imprecise exception make it difficult to restart execution after handling

---

# **MULTIPLE-ISSUE PROCESSORS**



# ILP: Scalar Pipeline (baseline)



**5-stage pipeline => 5 different instructions overlapped**

# Getting higher performance...

- Scalar pipeline limited to  **$CPI_{ideal} = 1$** 
  - *It can never fetch and execute more than **one** instruction per clock: **single-issue***
  - *It can be even worst due **stalls** added to solve hazards*
- Goal: To reach higher performance => **more parallelism** must be extracted from the program. In other words...**multiple-issue**
  - *It means fetching and executing more than **one** instruction per clock*
- Instruction dependences must be detected and solved: instructions must be *re-ordered* (**scheduled**) to achieve the highest ILP given the available resources.

# Getting higher performance...

---


- In a **multiple-issue pipelined** processor, we want to fetch and execute more than 1 instruction per clock

$$IPC_{ideal} > 1 \Rightarrow CPI_{ideal} < 1$$

- Example: **Dual-issue pipeline processor**  
*Ideal case:* max throughput would be to complete 2 Instructions Per Clock:

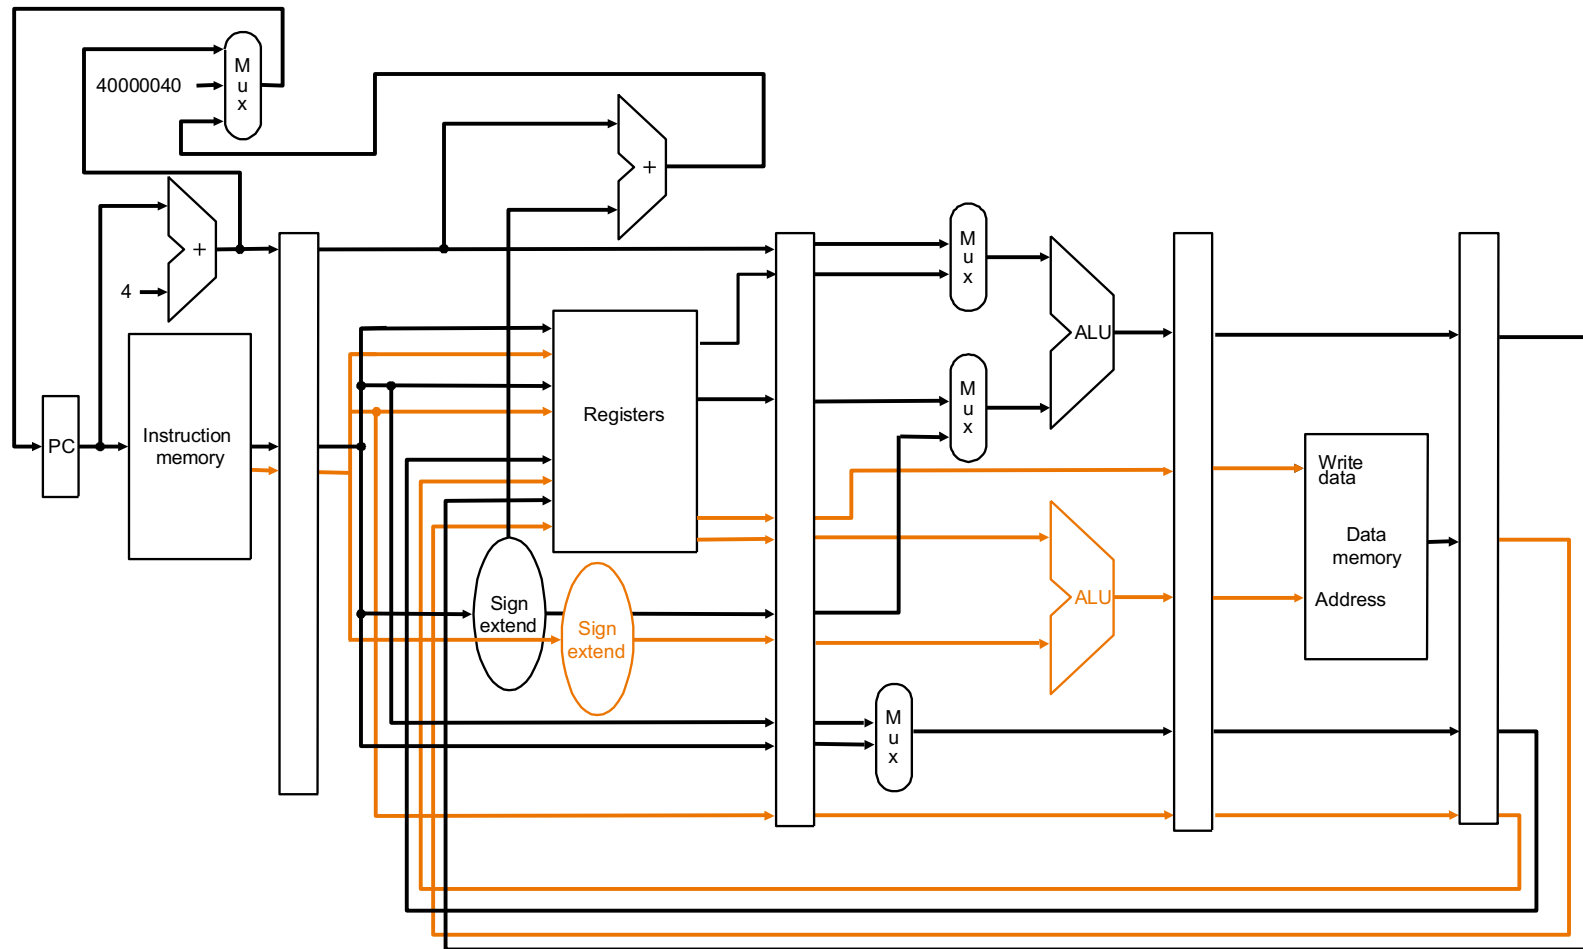
$$IPC_{ideal} = 2 \Rightarrow CPI_{ideal} = 0.5$$

# ILP: Dual-Issue Pipeline



**2-issue 5-stage pipeline => 2 x 5 different instructions overlapped**

# Dual-issue MIPS Pipeline Architecture



**2-instructions issued per clock:**


- 1 ALU or BR instruction
- 1 load/store instruction

**Multi-port RF: 4 read ports & 2 write ports**

# Recap


---

- *How to exploit the potential parallelism of the execution among independent instructions?*




Sequential (non pipelined) → IDEAL CPI > 1

# Several steps towards exploiting more ILP



# Several steps towards exploiting more ILP

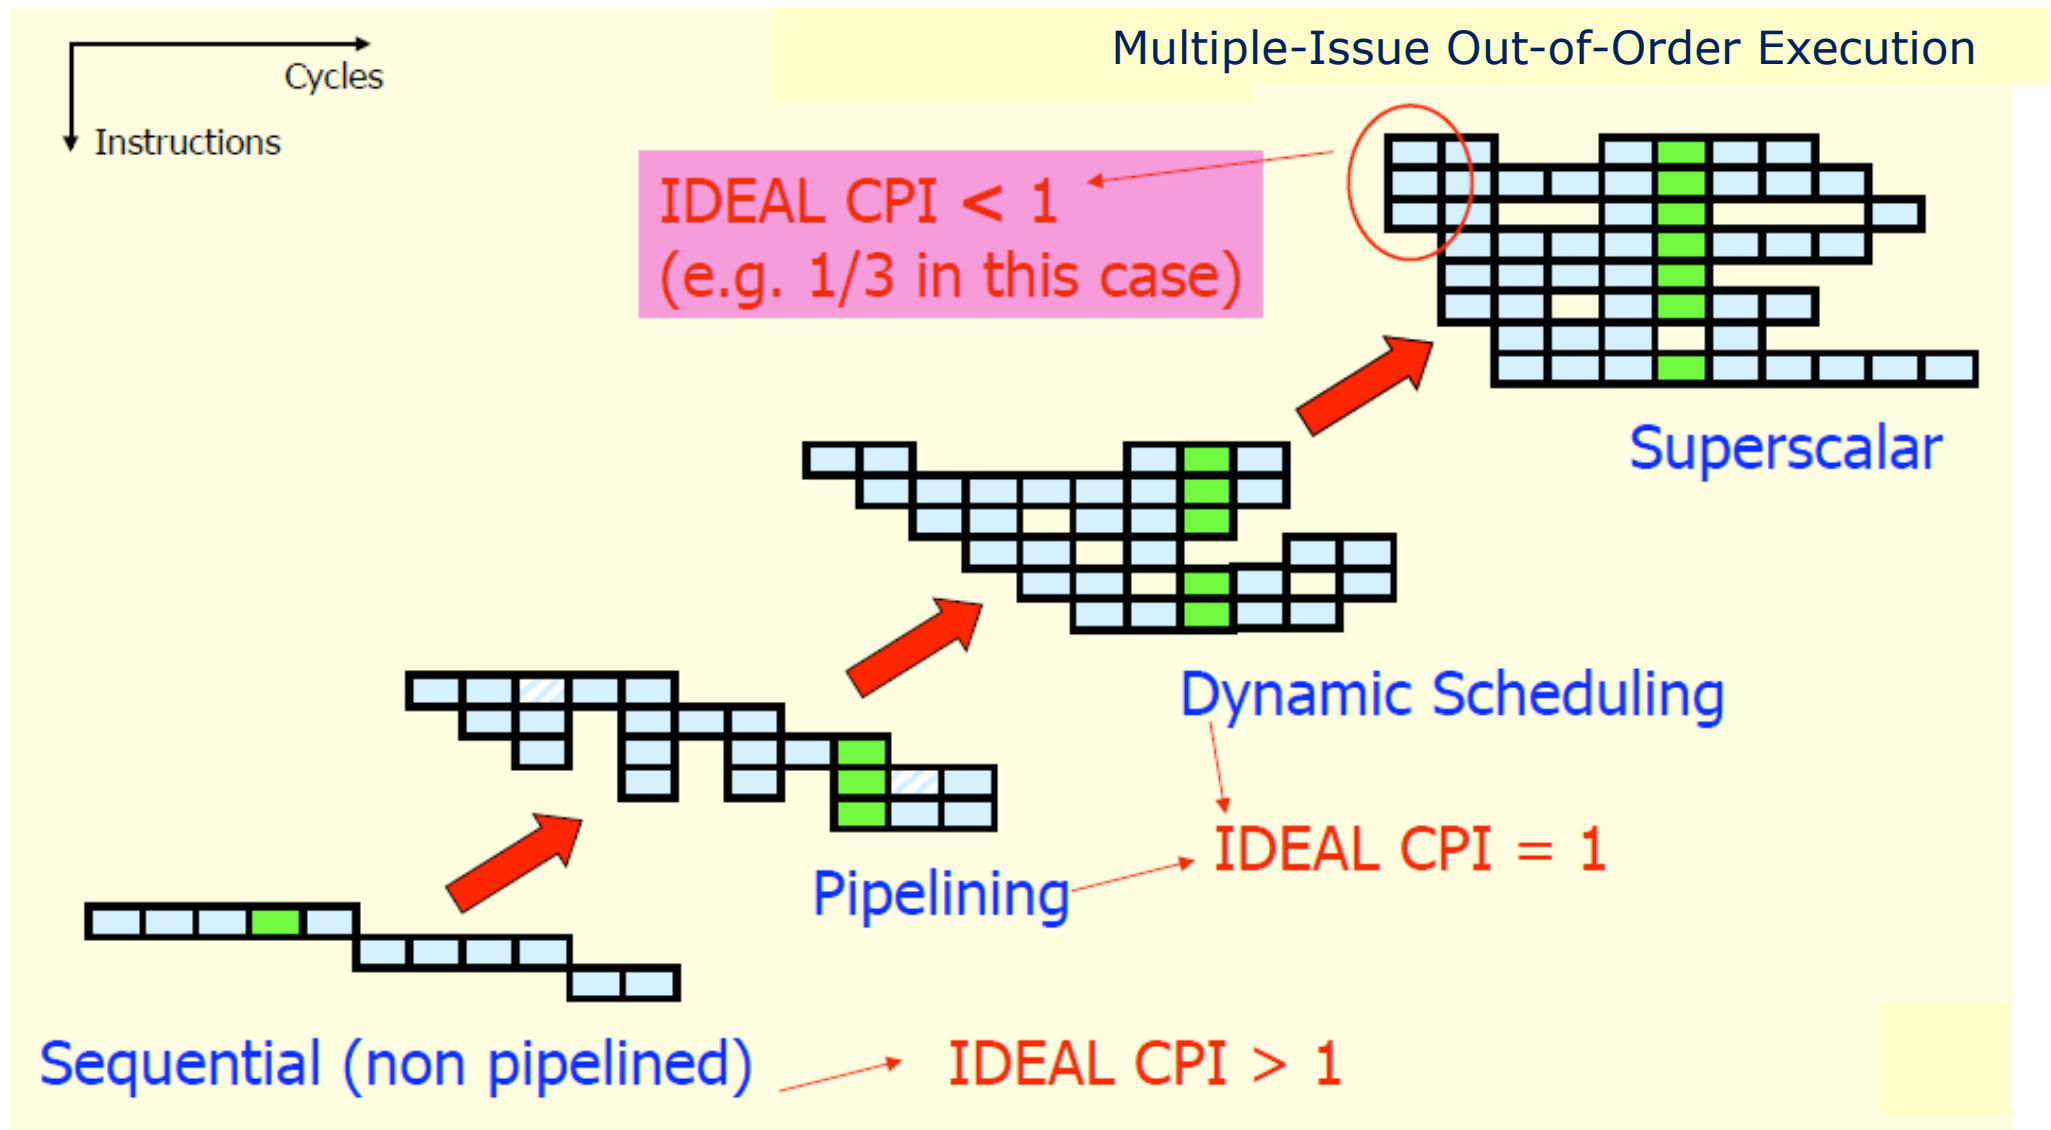





## Several steps towards exploiting more ILP

- **Dynamic Scheduling** can be applied to single-issue scalar processors but also to ***multi-issue processors***
  - => *Dual-issue MIPS processor and more in general **superscalar** processors*

# Several steps towards exploiting more ILP

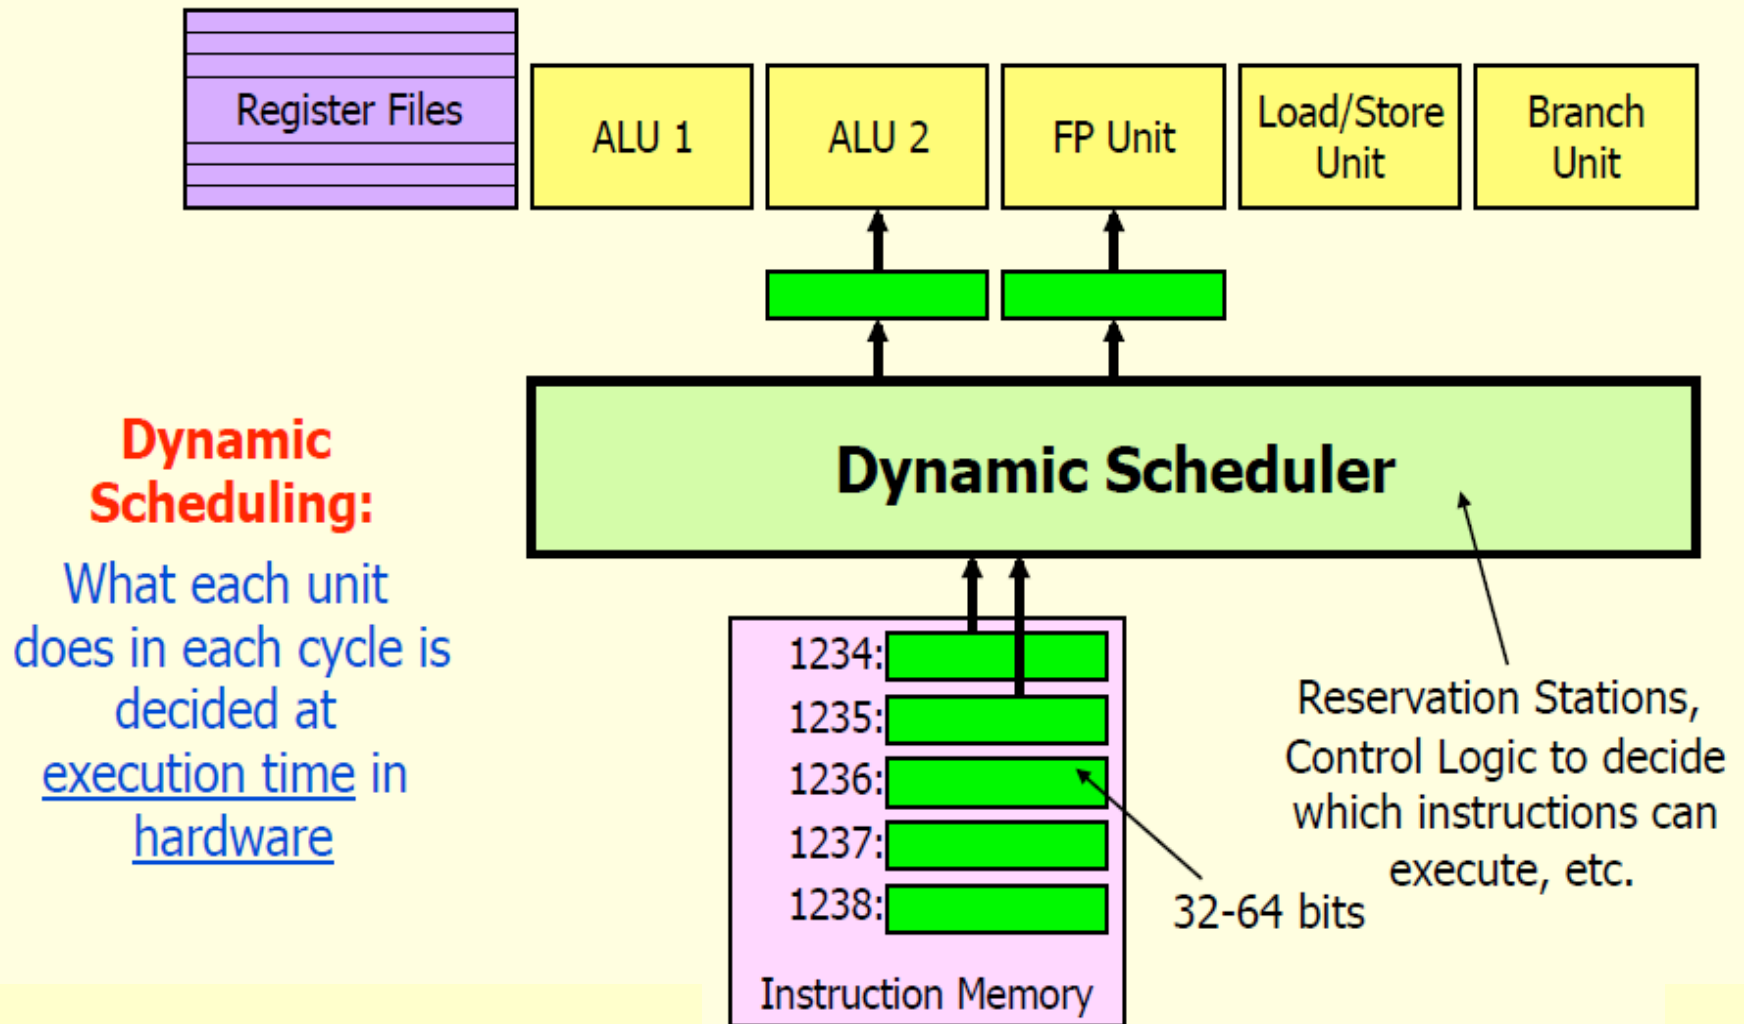


# Superscalar Execution



- Combining two main ideas:
  1. Why not fetching and beginning execution of more than one instruction per cycles?  
✓ **Multiple-issues**
  2. Why not to check and manage at runtime data and control dependencies?  
✓ **Dynamic scheduling:** Hardware reorder instructions execution so as to reduce stalls, maintaining data flow and exception behaviour.

# Dynamic Scheduler



# Recap on Dynamic Scheduling

- Instructions are *fetch*ed and *issued* in program order (**in-order-issue**)
- Execution begins as soon as **RAW hazards** are solved and operands are available  
=> *possible generation of **out of order execution*** –  
note: it is possible even with pipelined scalar architectures due to multi-cycle latencies.
- Out-of order execution introduces possible **WAR and WAW data hazards**.
- Out-of order execution generates ***out of order commit***.


## Recap on superscalar processors: Pros and Cons

- Main advantage:
  - Very high performance:  **$CPI_{ideal} = 1 / \text{issue-width}$**
- Disadvantages
  - Very complex logic and area cost to check and manage dependencies at runtime, i.e. to decide which instructions can be issued at every clock cycle;
  - Cycle time limited by scheduling logic (dispatcher and associated dependency checking logic)
  - It does not scale well: almost impractical to make issue-width greater than 4 ...

# Why not using static scheduling?

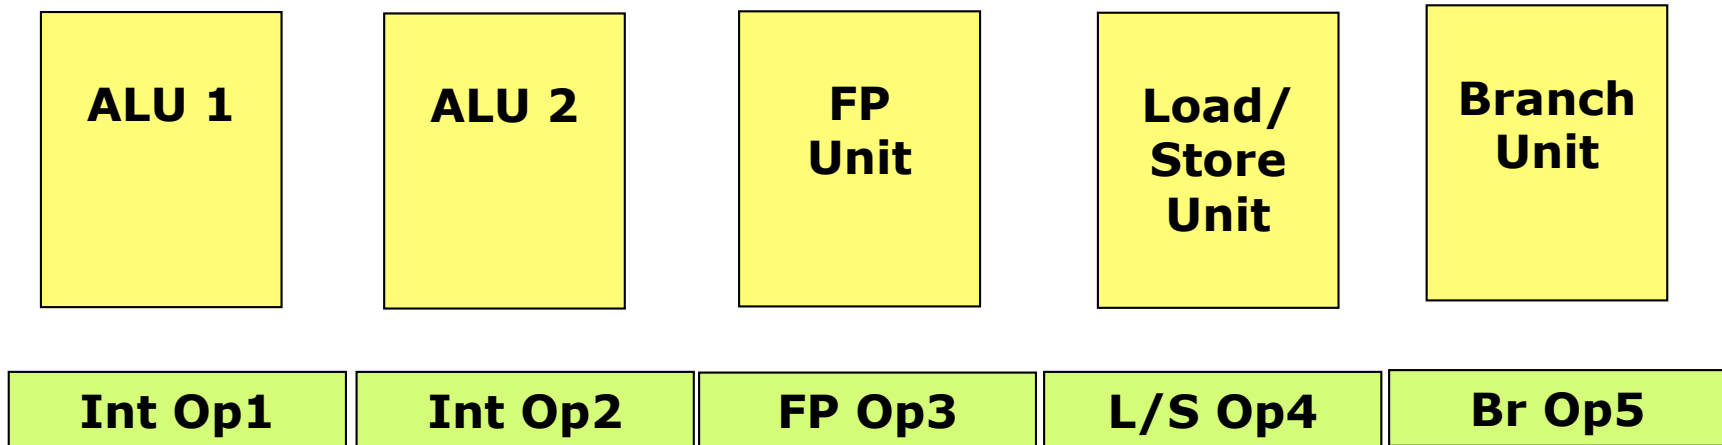
- Static detection and resolution of dependences
  - ⇒ **static scheduling** done by the **compiler**
  - ⇒ instruction dependences are avoided by code reordering at compile time.
- Output of the compiler: instructions are reordered into dependency-free parallel instructions, otherwise NOPs are introduced in the code.
- Typical architecture: **VLIW (Very Long Instruction Word)** processors expect **dependency-free code** generated at static time by the compiler.

# Very Long Instruction Word: An Alternative Way of Extracting ILP

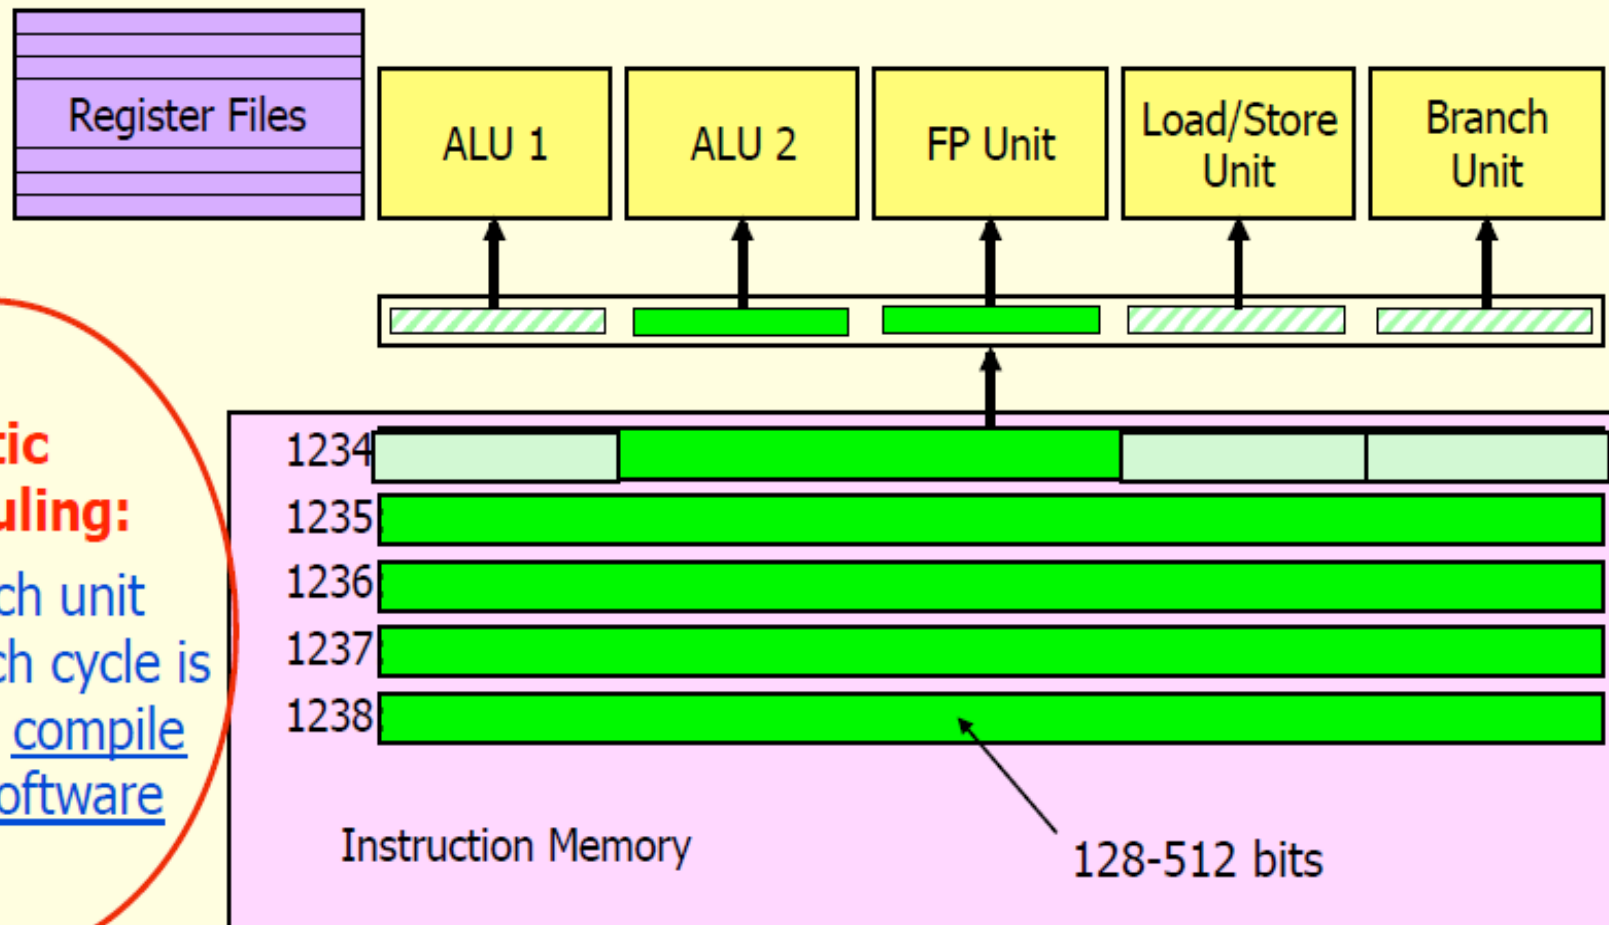




# VLIW: Very Long Instruction Word



# (Statically Scheduled) Very Long Instruction Word Processor (VLIW)



## Static Scheduling:

What each unit does in each cycle is decided at compile time in software

# Main Limits of Static Scheduling

---

- Unpredictable branch behavior: the code parallelization is limited to basic blocks.
- Unpredictable cache behavior: Variable memory latency for hits/misses.
- Complexity of compiler technology: The compiler needs to find a lot of parallelism in order to keep the multiple functional units of the processors busy
- Code size explosion due to insertion of NOPs
- Low code portability and binary code incompatibility: Consequence of the larger exposure of the microarchitecture (= implementation details) at the compiler in the generated code
- Low performance portability

# Recap on Static vs Dynamic Scheduling

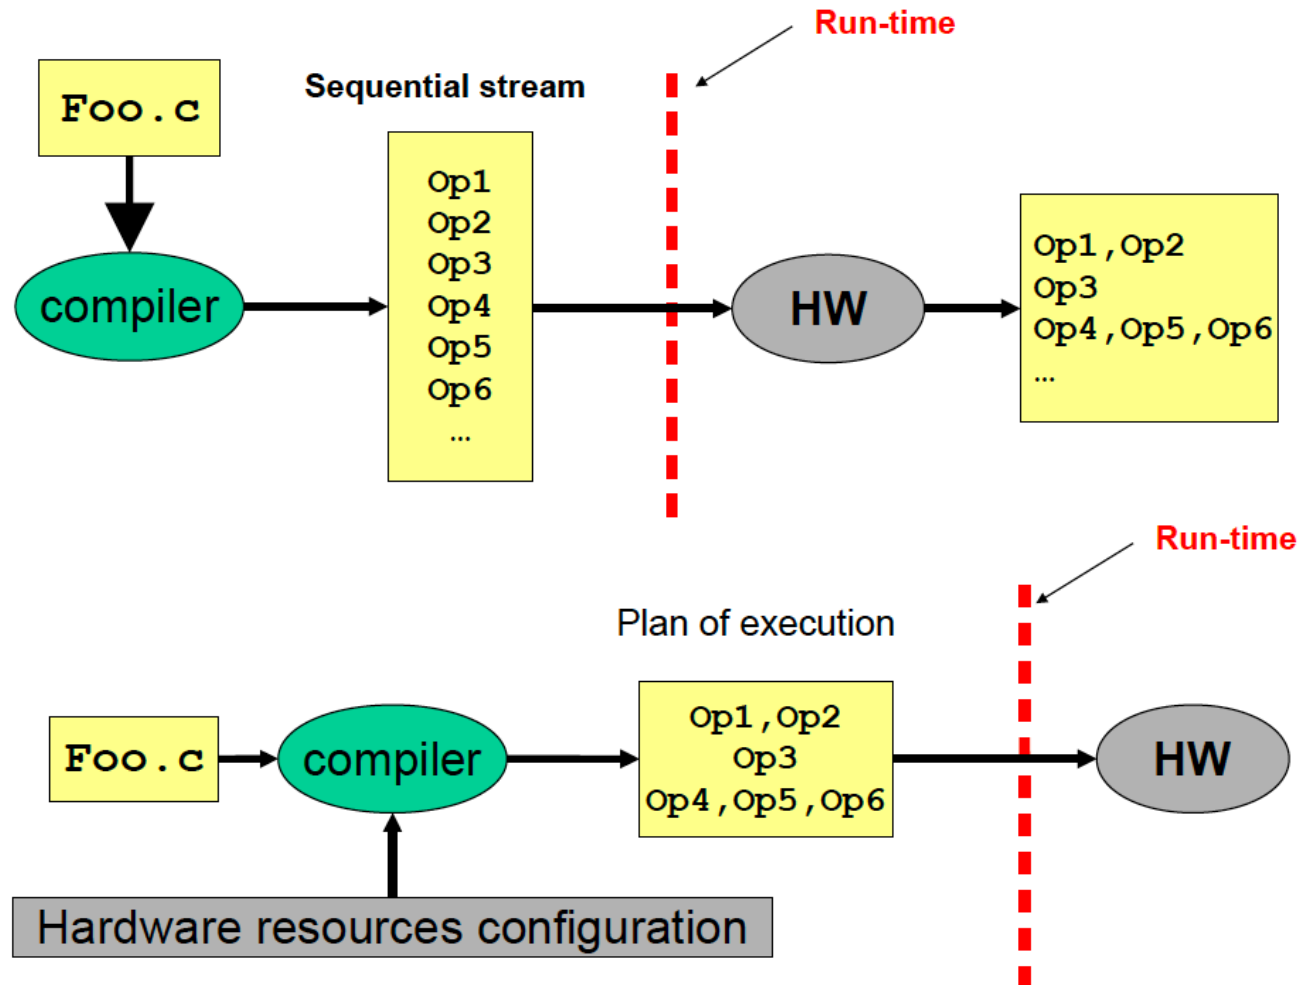
- Two strategies to support ILP:
  - **Static Scheduling:** Rely on the compiler for identifying potential parallelism
  - **Dynamic Scheduling:** Depend on the hardware to locate parallelism

# ILP: Static vs Dynamic Scheduling

- Instruction Scheduling: Deciding **when** and **where** to execute an instruction
  - *i.e.* in which cycle and in which functional unit
    1. For a superscalar processor it is decided **at runtime**, by the hardware logic.
    2. For a VLIW processor it is decided **at compile time**, by the compiler, and therefore by a SW program
      - Suitable for embedded processors: Simpler HW design (no dynamic scheduler), smaller area and power consumption ... and cheap

# ILP: Static vs Dynamic Scheduling

- Who decides about the ILP?*



# Issue-width limited in practice

- The **issue width** is the number of instructions that can be issued in a single cycle by a **multiple issue processor**
- When superscalars were invented, 2- and rapidly 4-issue width processors were created.
- Due to the intrinsic level of parallelism of a program, it is hard to decide which 8, or 16, instructions can execute every cycle!

## How to overcome the ILP limitations?

### ***Solution: Introduce more levels of parallelism***

- Multi-threading (Hyperthreading)
- Multi-processing and multi-cores
- Data-level Parallelism: Vector Processors and GPUs
- Heterogeneity: Host processor and co-processor